ELTE | FACULTY OF INFORMATICS

Local Code Assistants, Agents, Coding Agents

Continue Dev

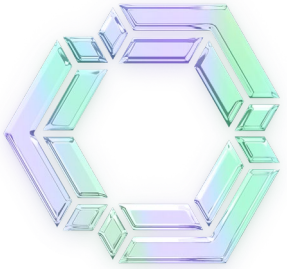

- „Local” Copilot-like extension
- Pros:
  - Privacy
  - Cost(?)
  - Customizable models
- Cons:
  - Speed(?)
  - Accuracy

Continue Dev

- Install from extensions
- Use continue tab
  - Try with OpenAI API key
- Functionality similar to Copilot Chat-mode
- Custom commands/prompts available

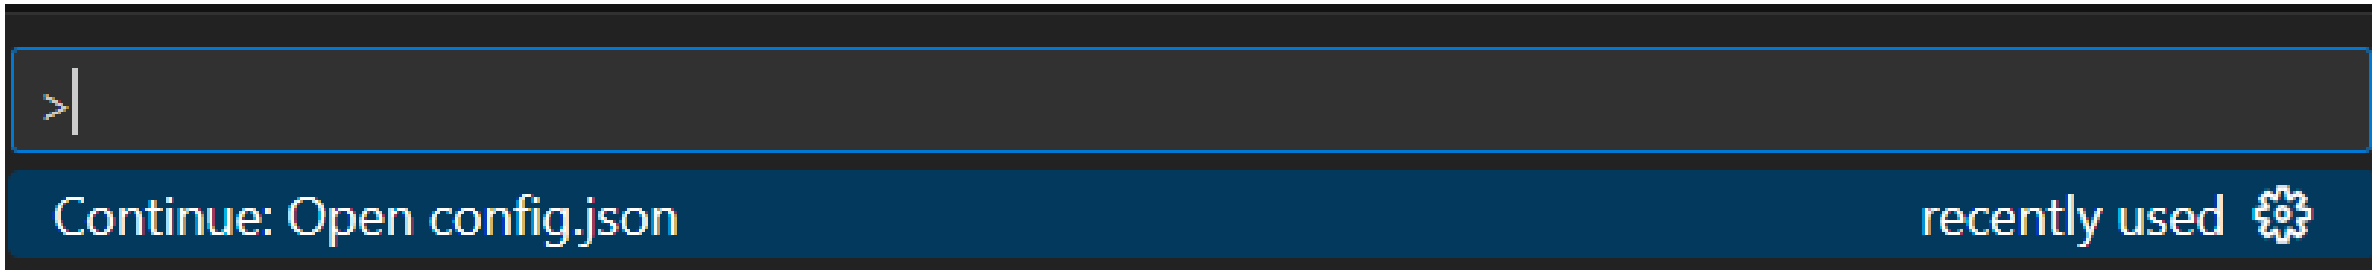

Continue Dev

- Local LLM-s with Huggingface
- https://huggingface.co/docs/text-generation-inference - Huggingface TGI
- Models at: https://huggingface.co/
- Supported models: https://docs.continue.dev/setup/select-model

Prompting chat models

- System prompt
  - General definition for the whole conversation, does not get cut out of history
- User-Assistant prompts form the history which is limited. Breaking alternation is often allowed.
- Textual tool outputs might also be included in the conversation

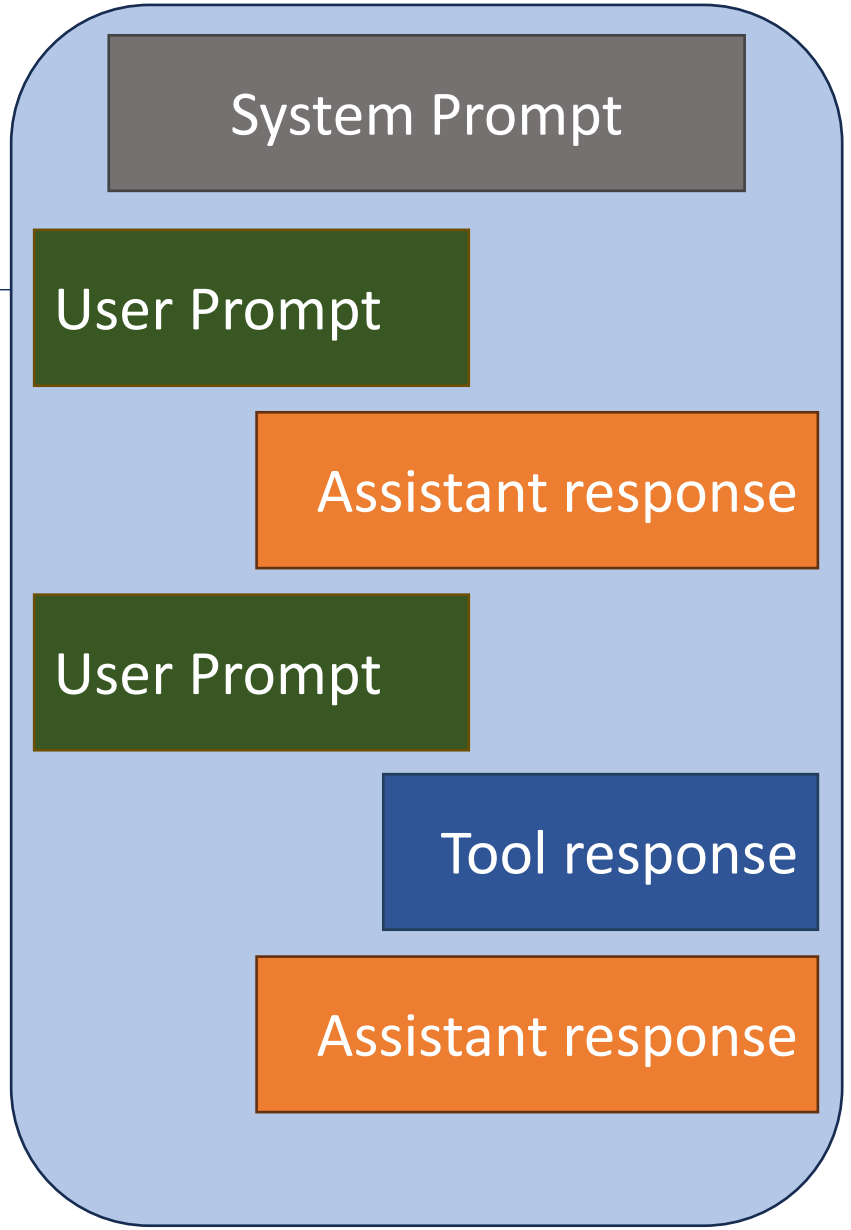

In-Context Learning

- Few-shot prompts include the task definition and a few examples.

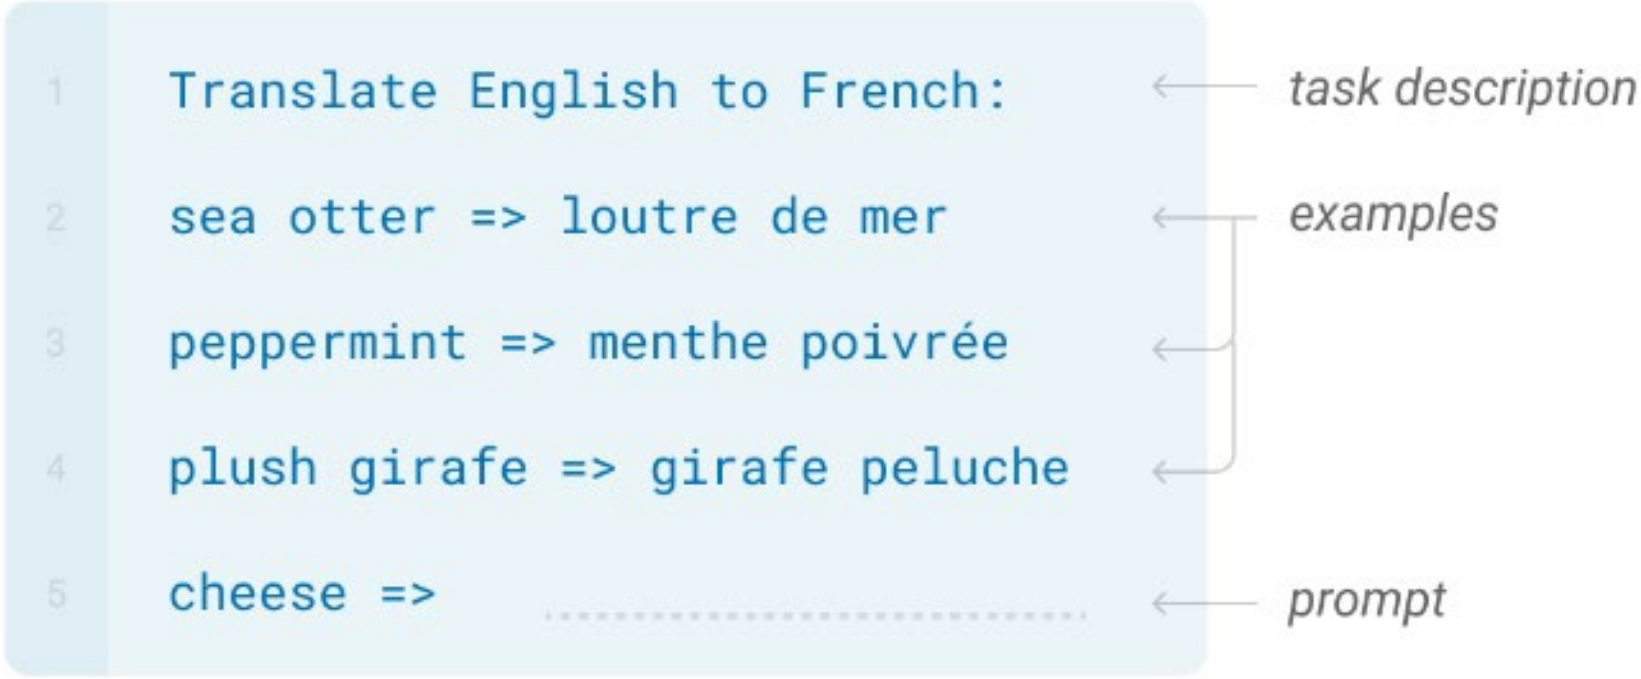

Source: Brown, et. al. 2020.

Problem Decomposition

- LLMs have limited single-step reasoning capabilities
- They are good at small common sense reasoning steps
- Hard problems should be decomposed into a Chain of Thoughts (CoT), where each thought is just a small step towards the solution.
- Latest models have built-in (by alignment) CoT capabilities

CoT Example

Which one is older Bosch or the Eiffel tower?

The Eiffel Tower is older than Bosch. The Eiffel Tower was completed in 1889, while Bosch was founded in 1886.

Which one is older Bosch or the Eiffel tower?
Give your answer step-by-step!

1. The Eiffel Tower was completed in 1889.
2. Bosch is a company that was founded in 1886.
3. Therefore, Bosch is older than the Eiffel Tower by 3 years.

Provide external knowledge

- Specific knowledge is missing from LLMs
- Models frequently hallucinate

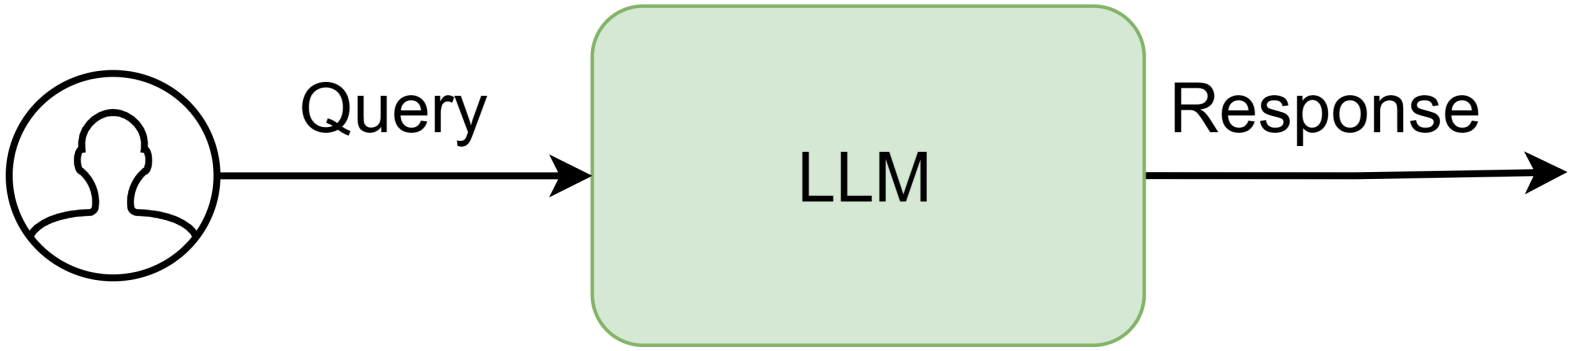

Retrieval Augmented Generation (RAG)

- Using embedding models to encode the meaning of text chunks (1-3k dimensional vectors)
  - Usually encoder-style LLMs finetuned for similarity search
- VectorDB – similarity search
- Hybrid search possible
- Adding knowledge from the same topic to the prompt

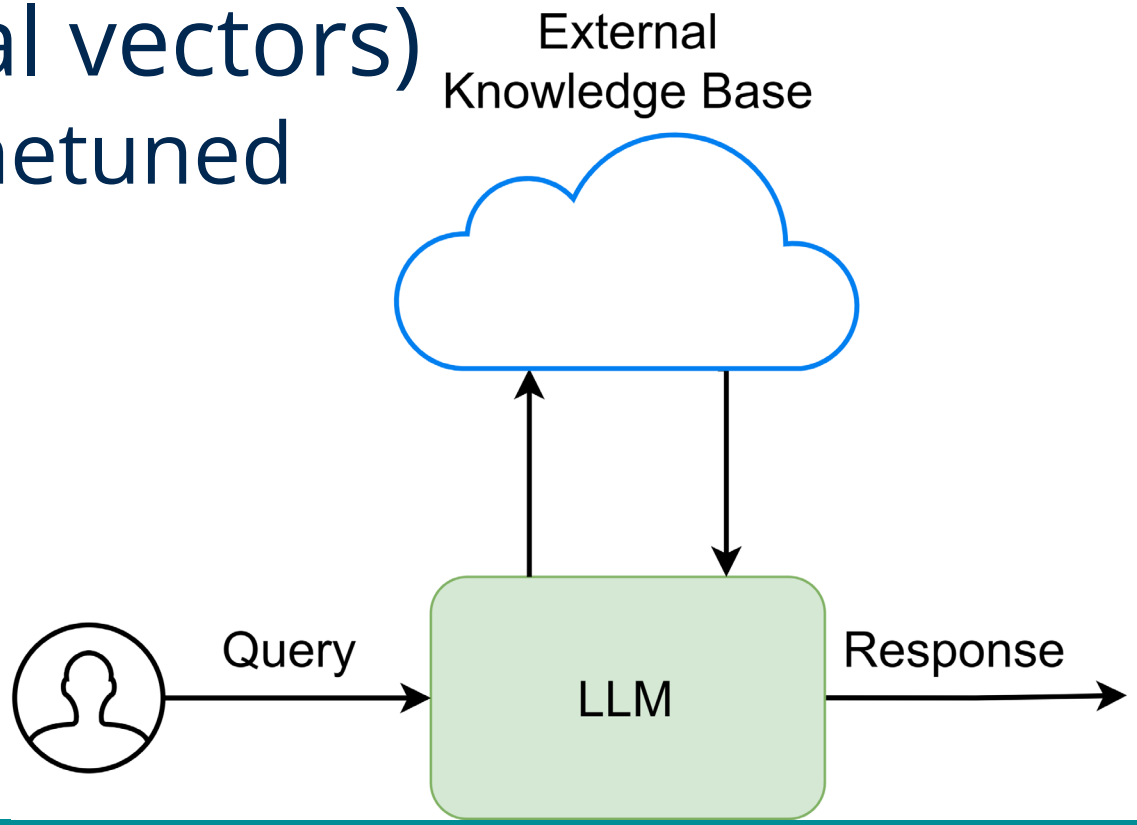

2-Step RAG

- Questions are not always similar to the relevant knowledge:
  - References in the chat
  - Background information is better than specific
- Answers are assembled in a separate step

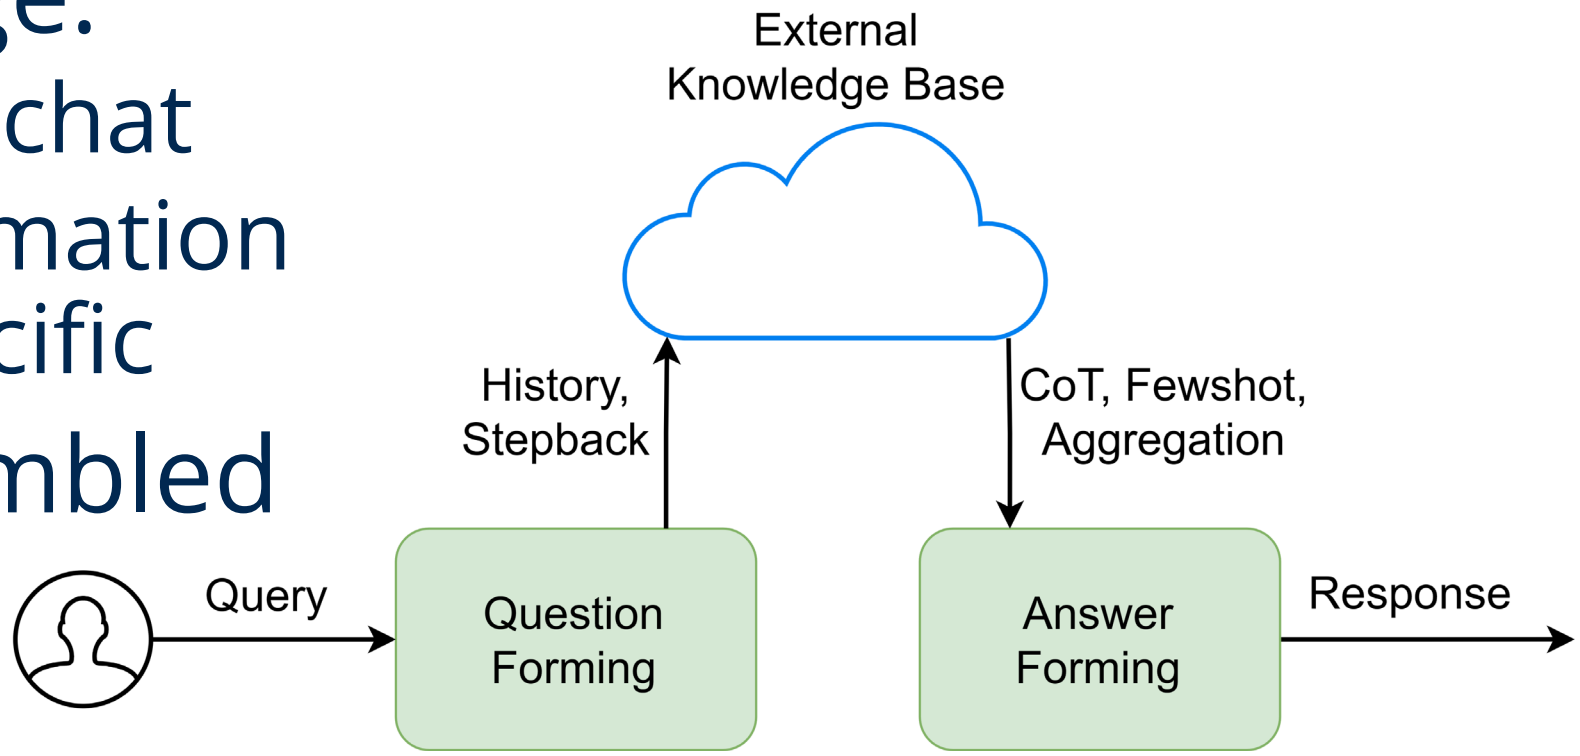

Reflection

- An extra reflection node (which reflects on the suitability of the output) can filter a large chunk of remaining hallucinations

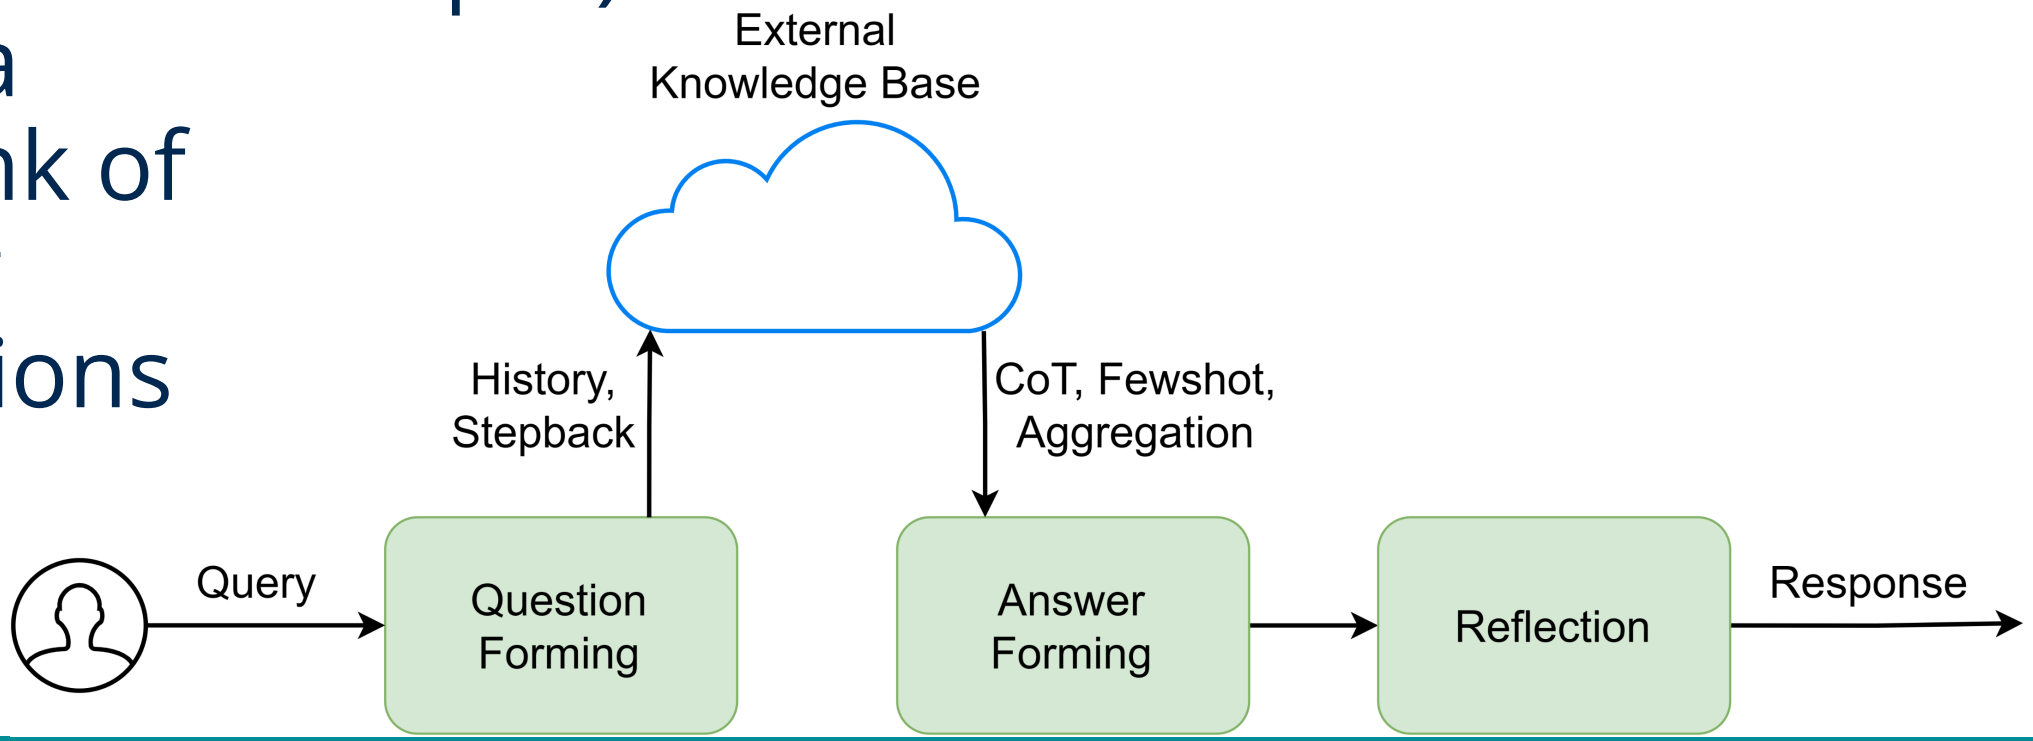

Tooling

- Tools that can be called via text parameters and return text input are suitable for LLM chains
  - Programming kernel
  - Search engines
  - Knowledge graphs
  - Action handlers
  - Application APIs
- JSON is a popular format

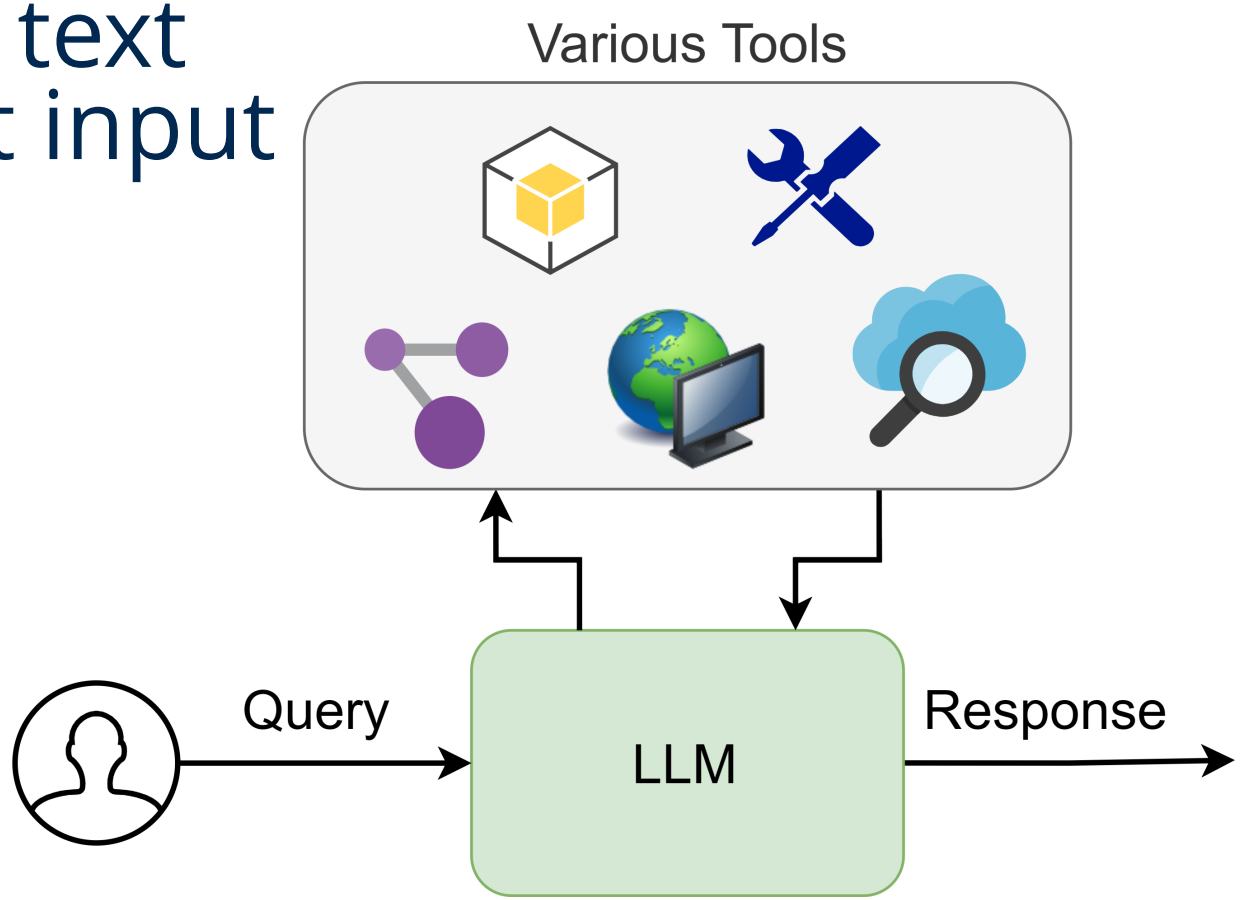

Agents

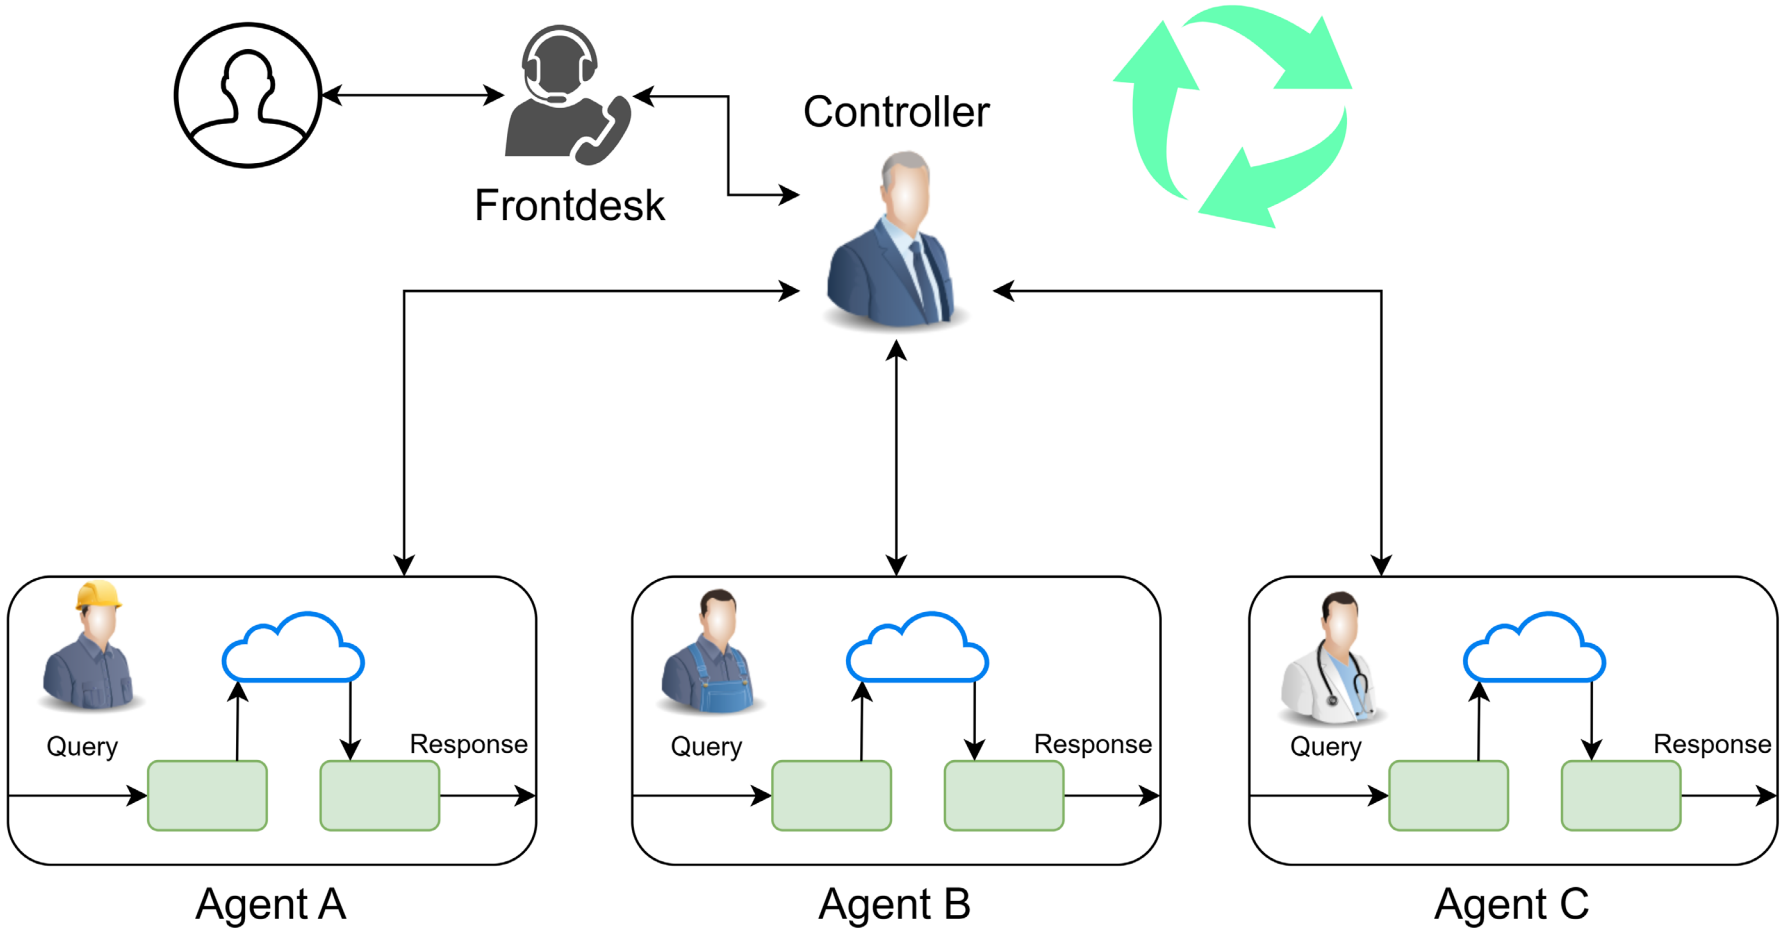
Frontdesk
Controller
Query
Response
Agent A
Query
Response
Agent B
Query
Response
Agent C

Agents

- Specialist Agents are LLMs with external knowledge/tools that specialize in performing a single task
- Controller is an LLM, that uses a loop of planning, task execution and evaluation

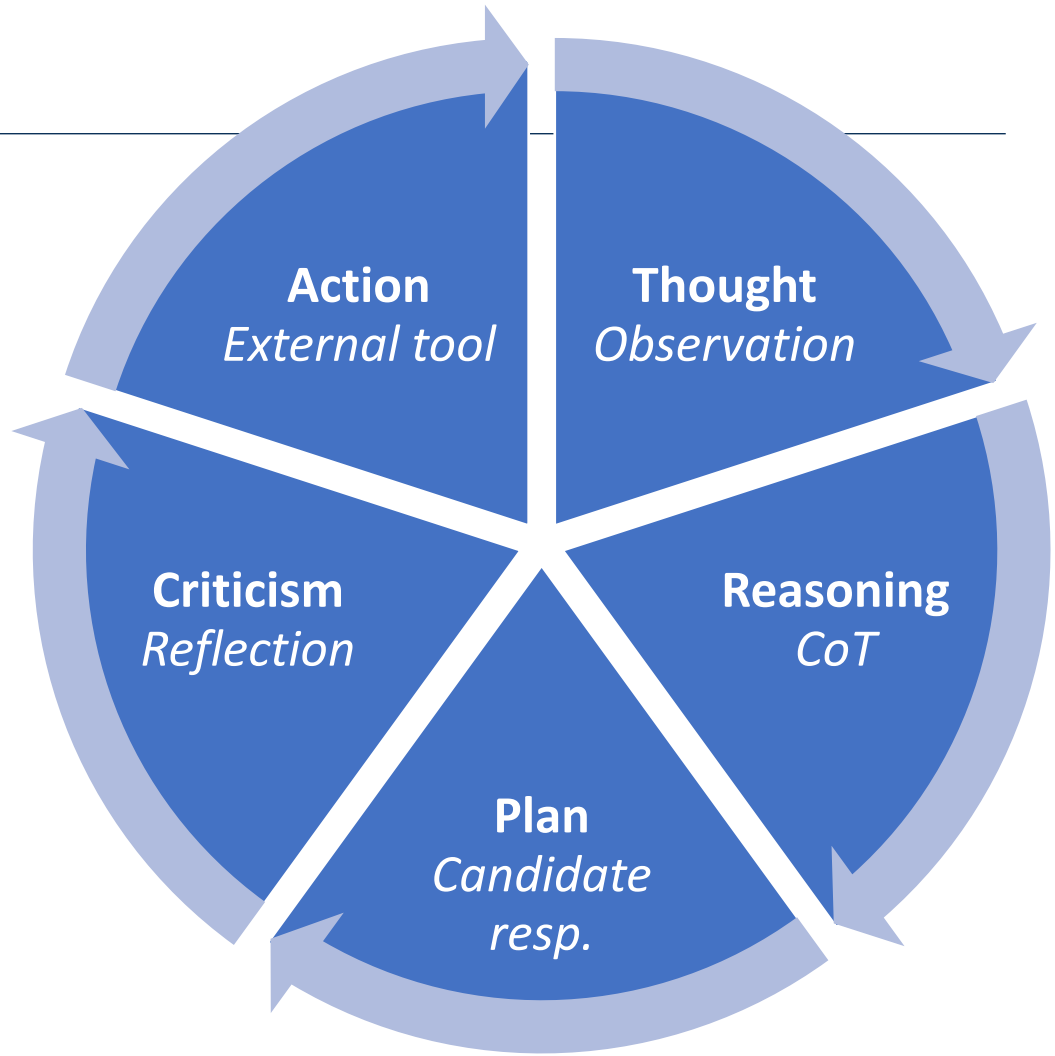

Example: Design Assistant

- Integration of scientific tools and LLM pipelines
- Immature stage: 1st conference on LLM-aided design this June
- Agentified chat systems are viable

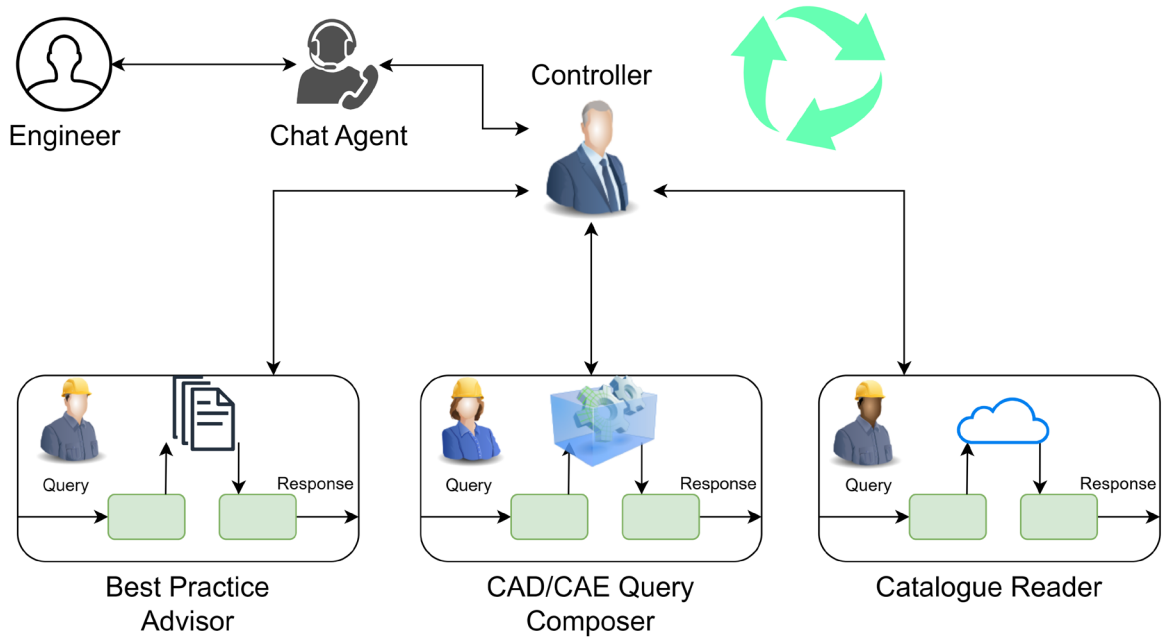

Example: Design Assistant

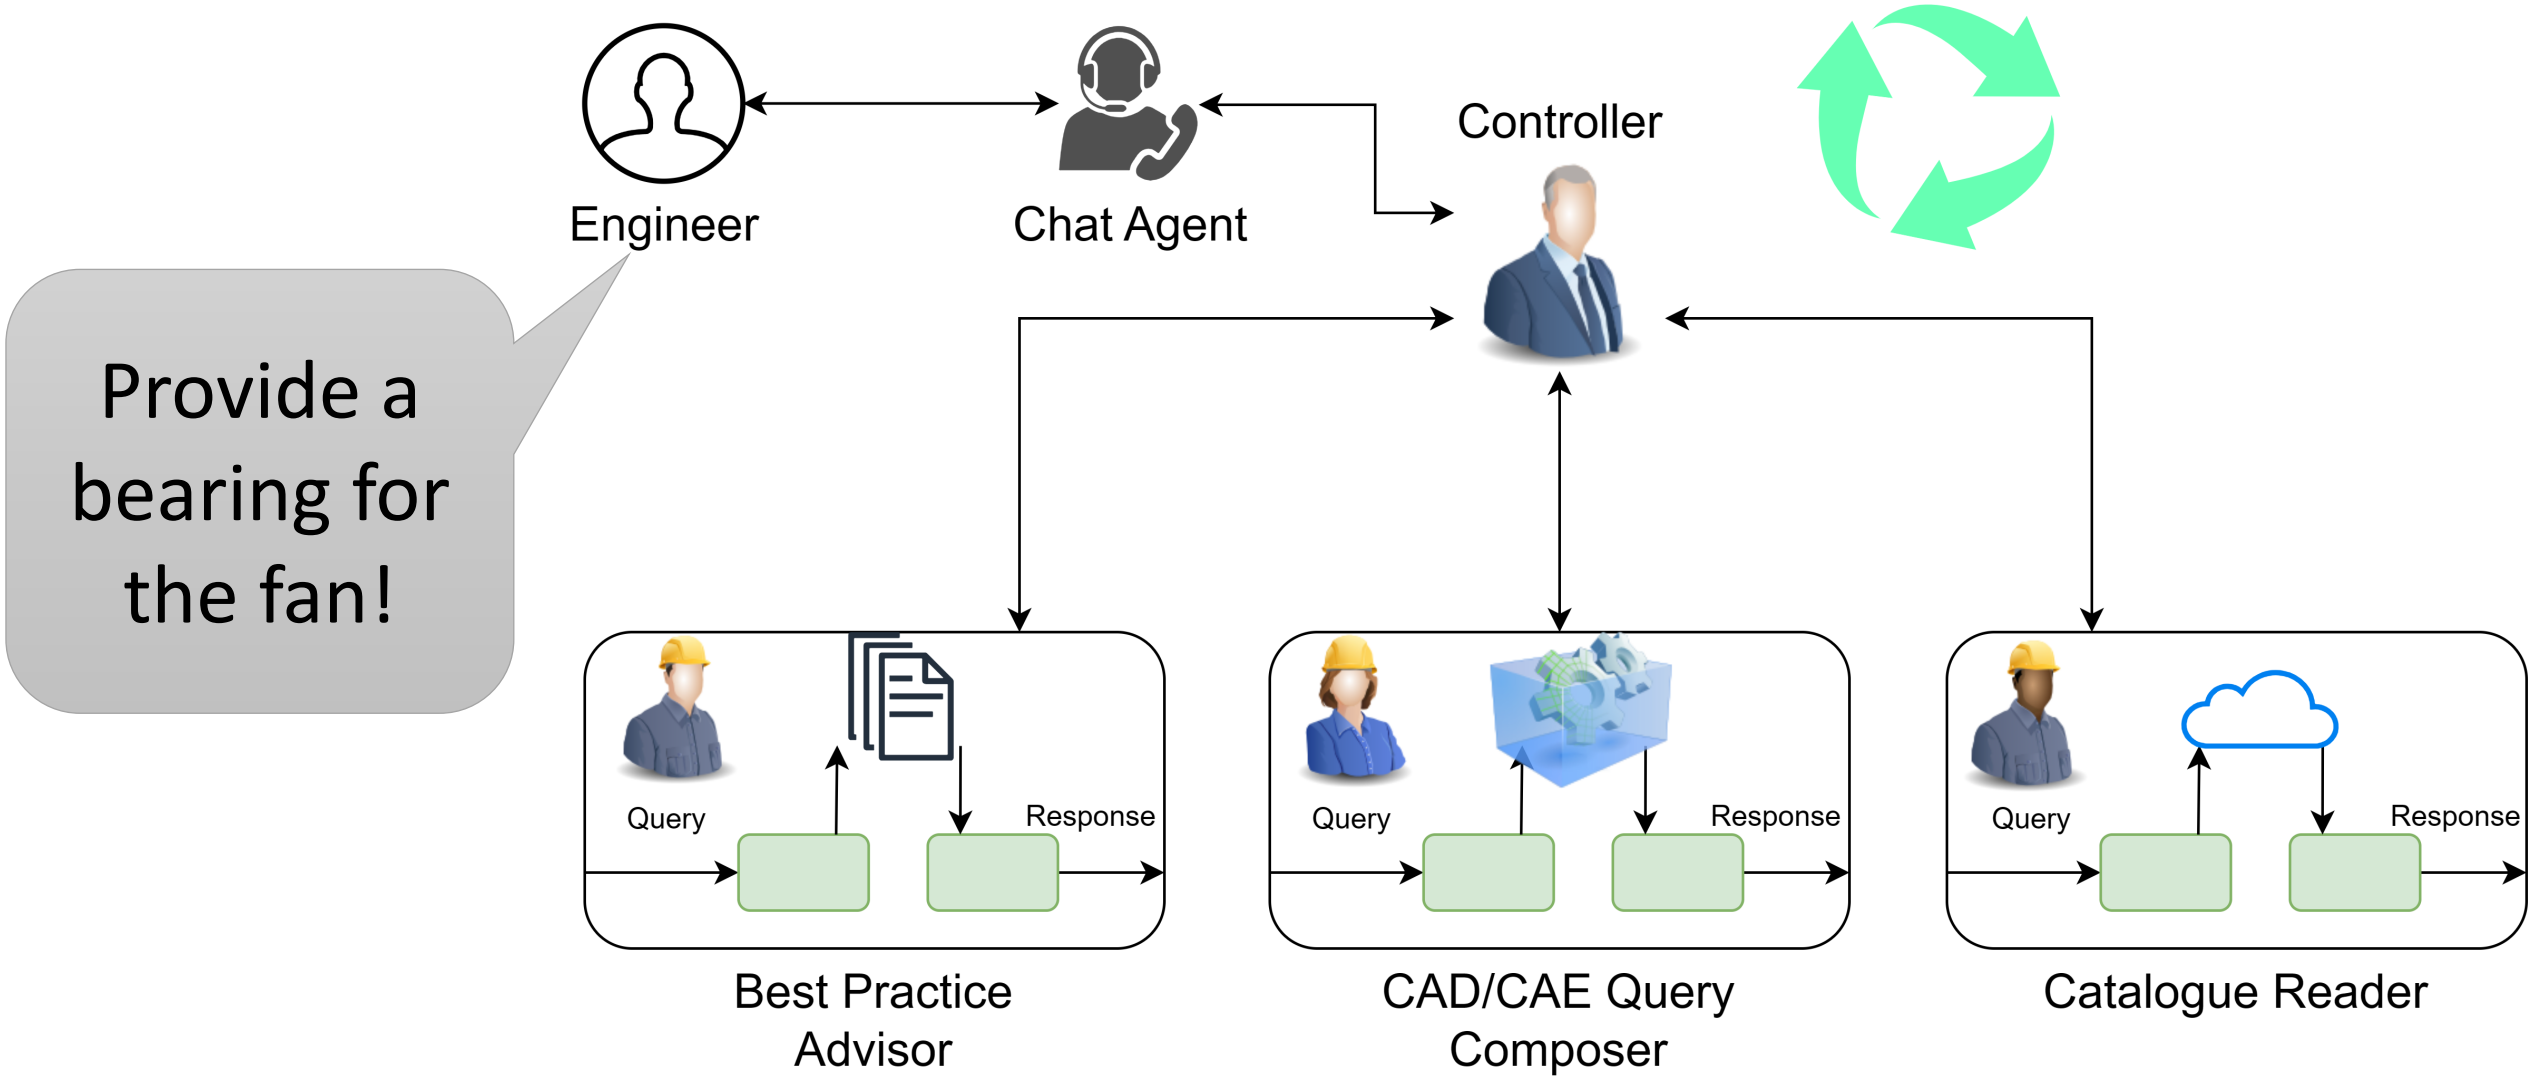

Example: Design Assistant

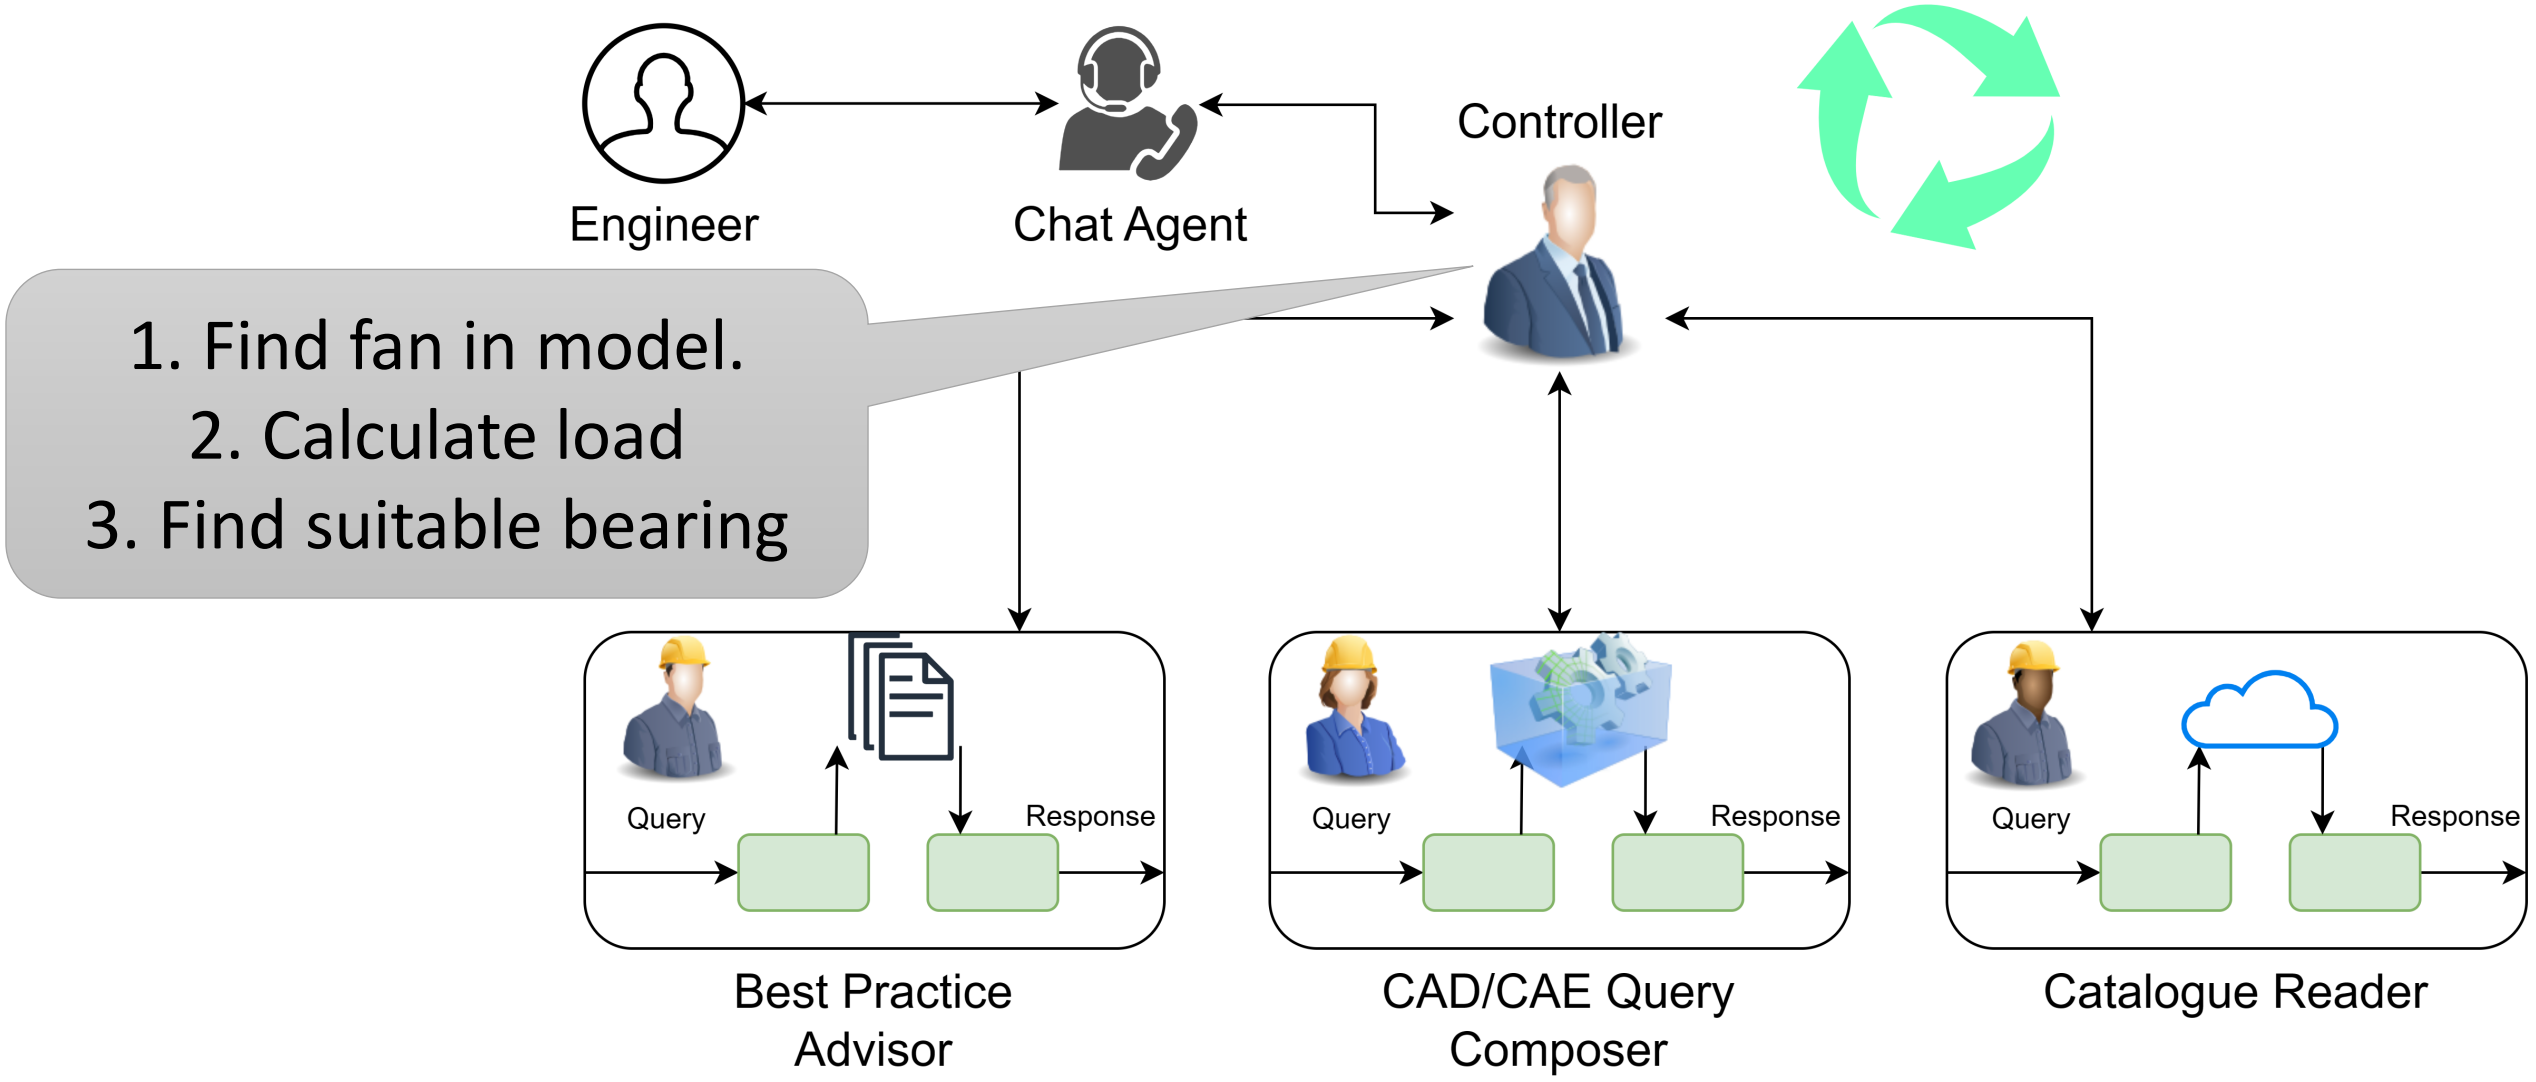

Example: Design Assistant

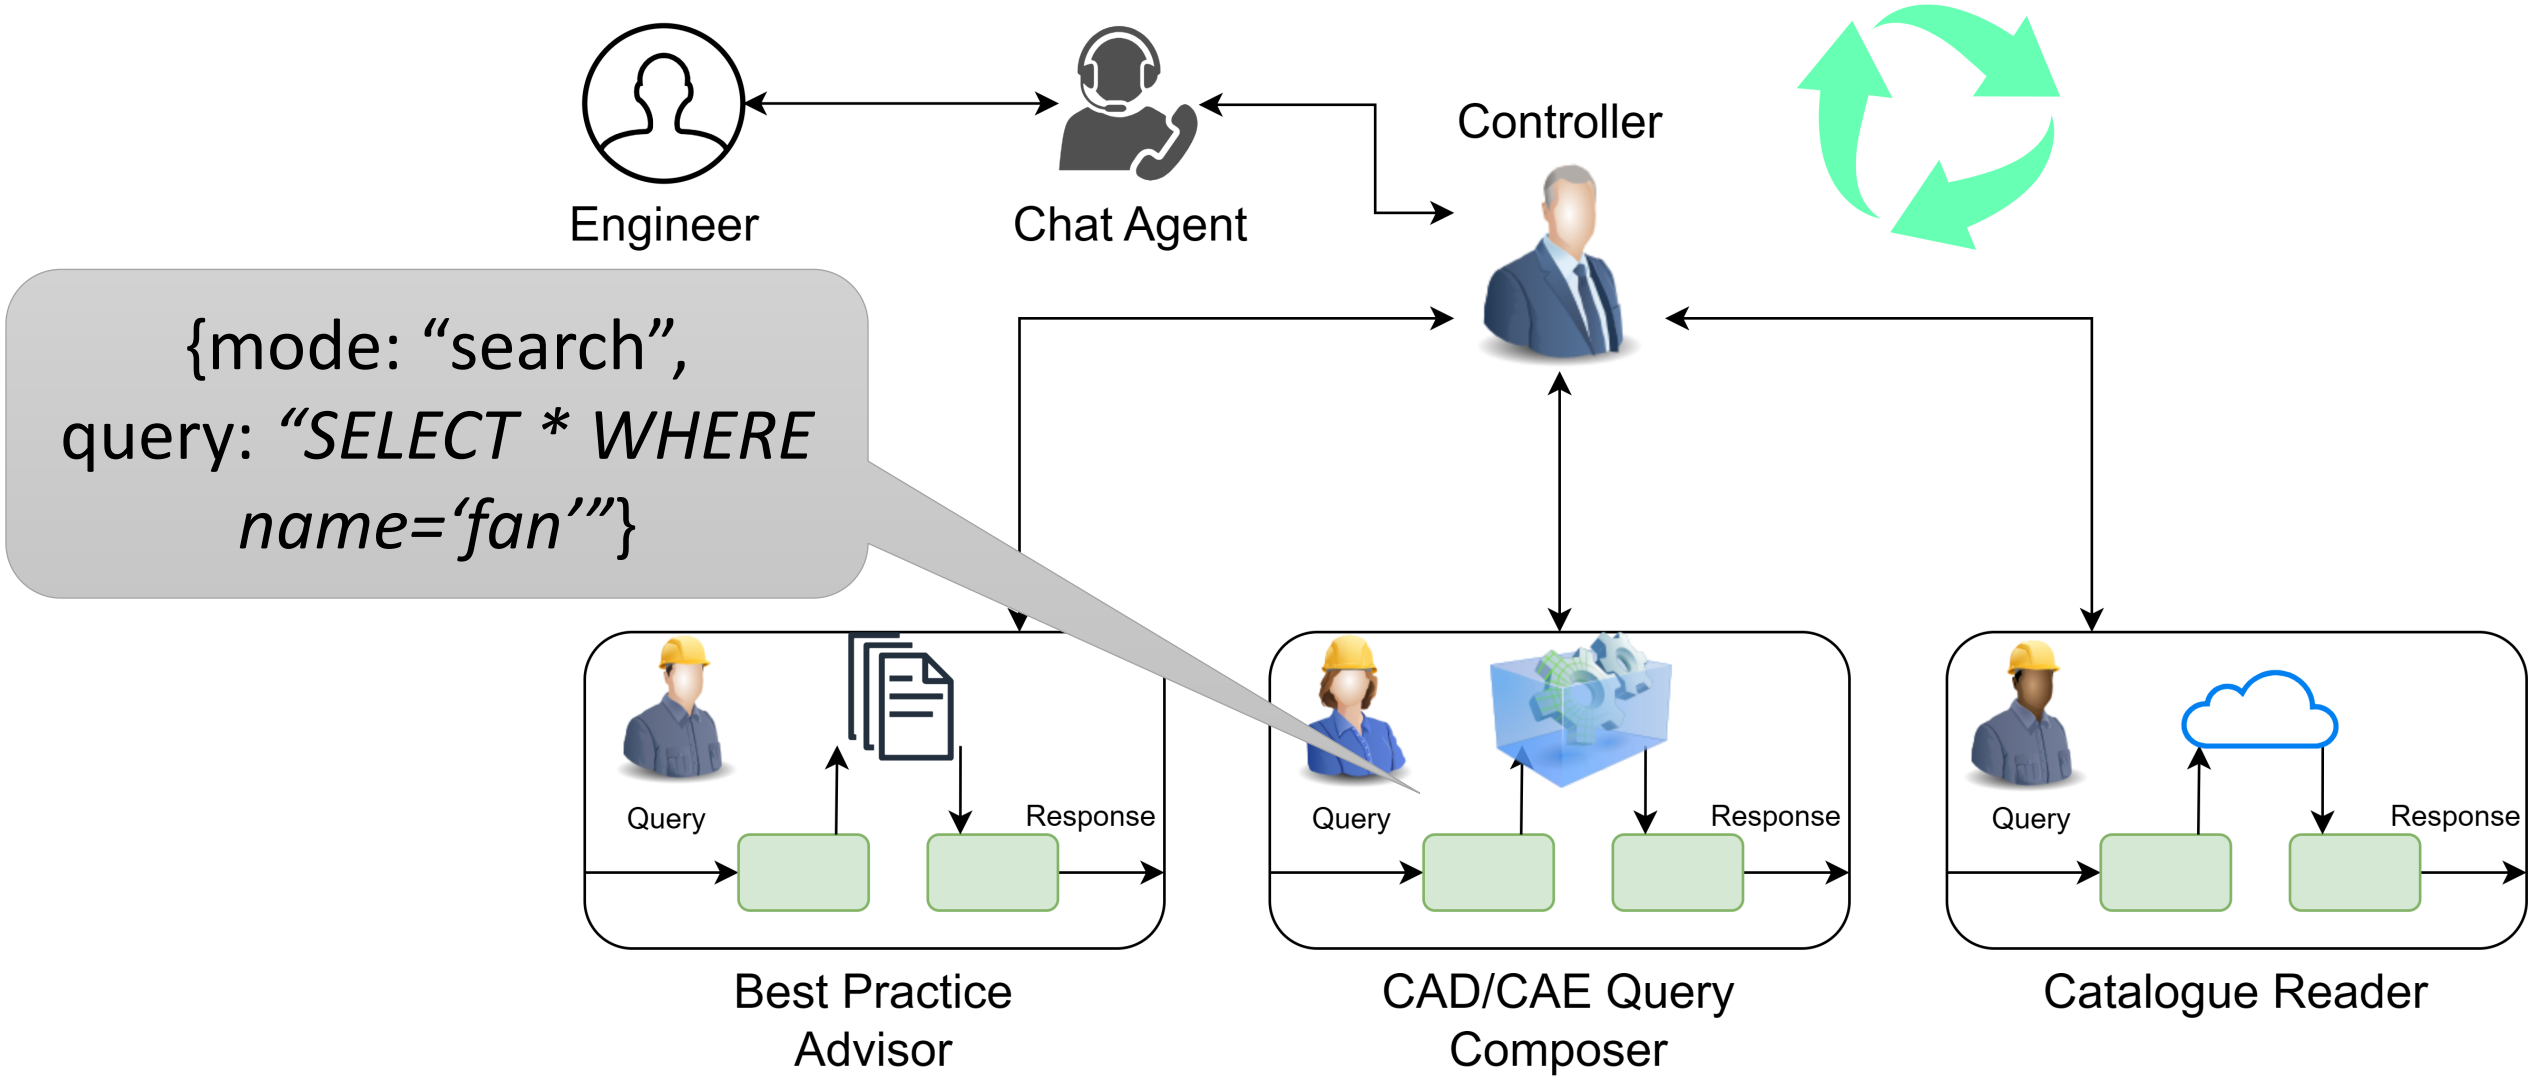

Example: Design Assistant

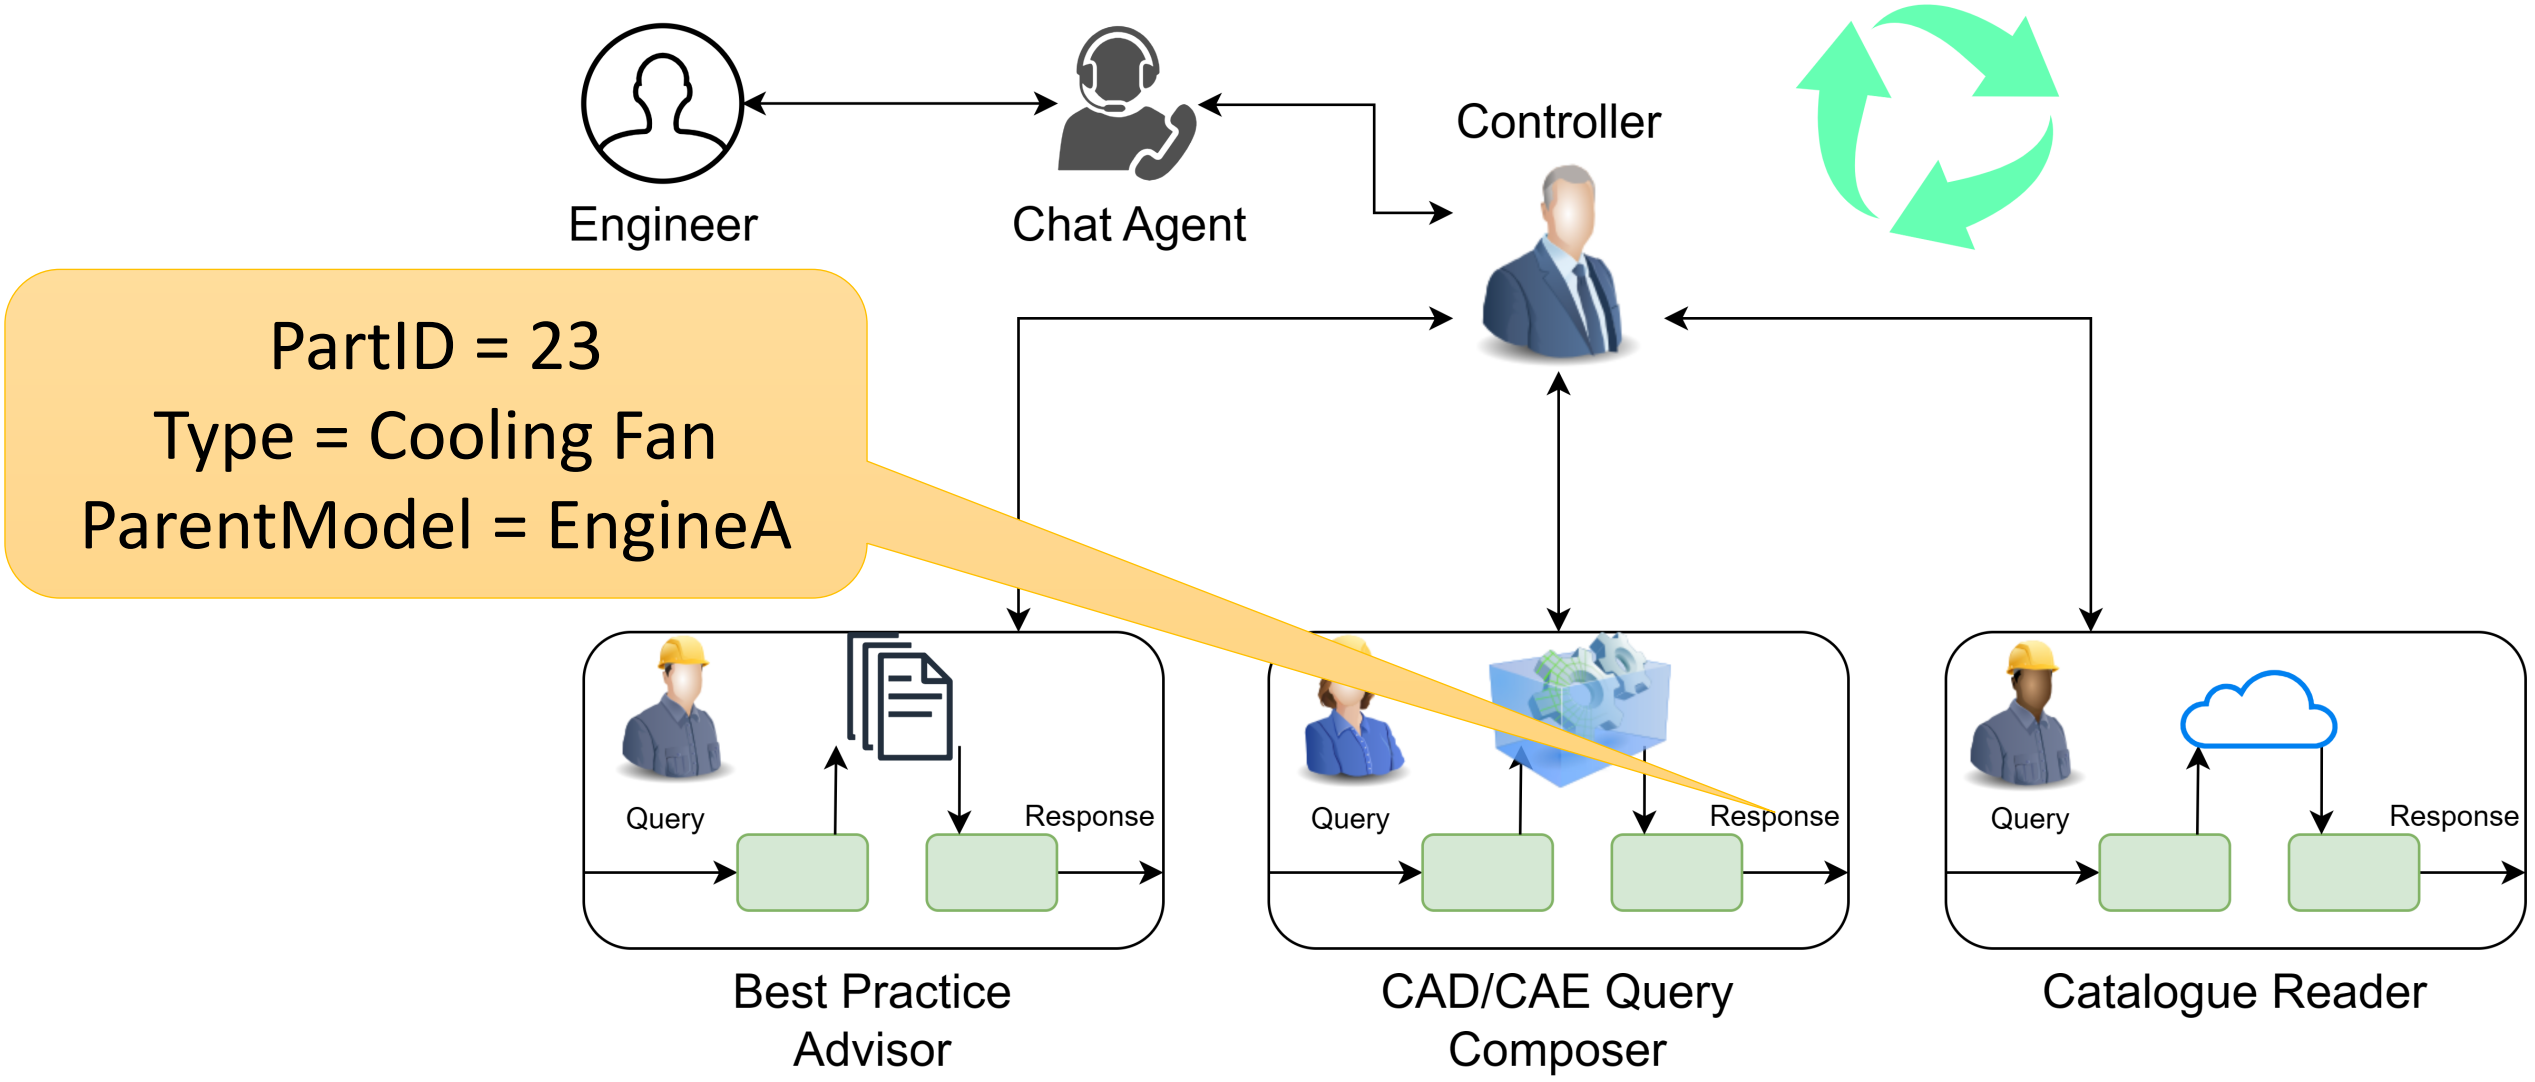

Example: Design Assistant

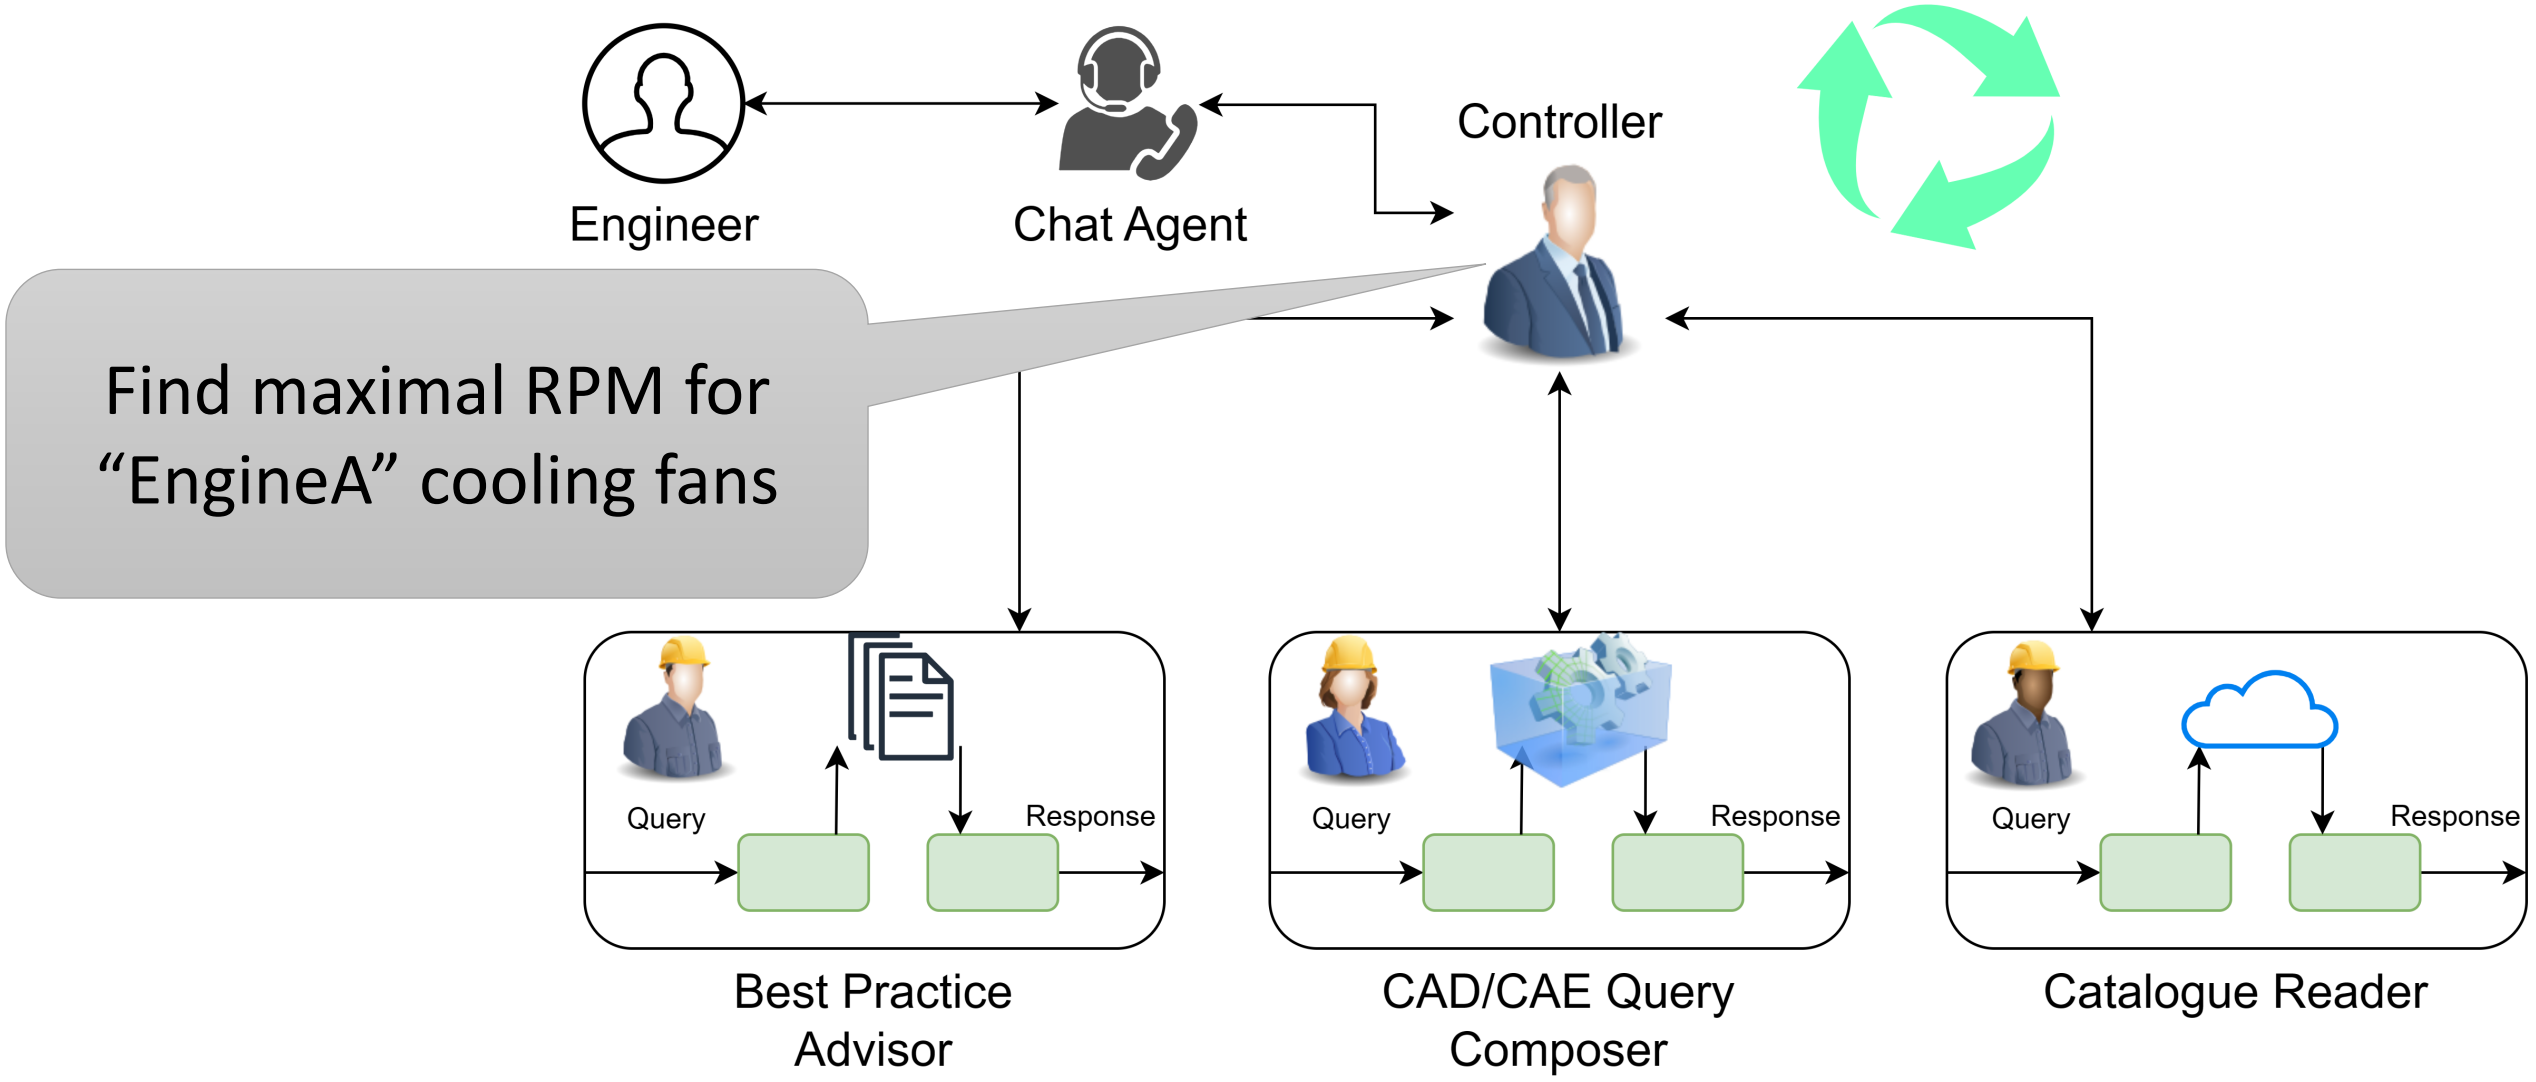

Example: Design Assistant

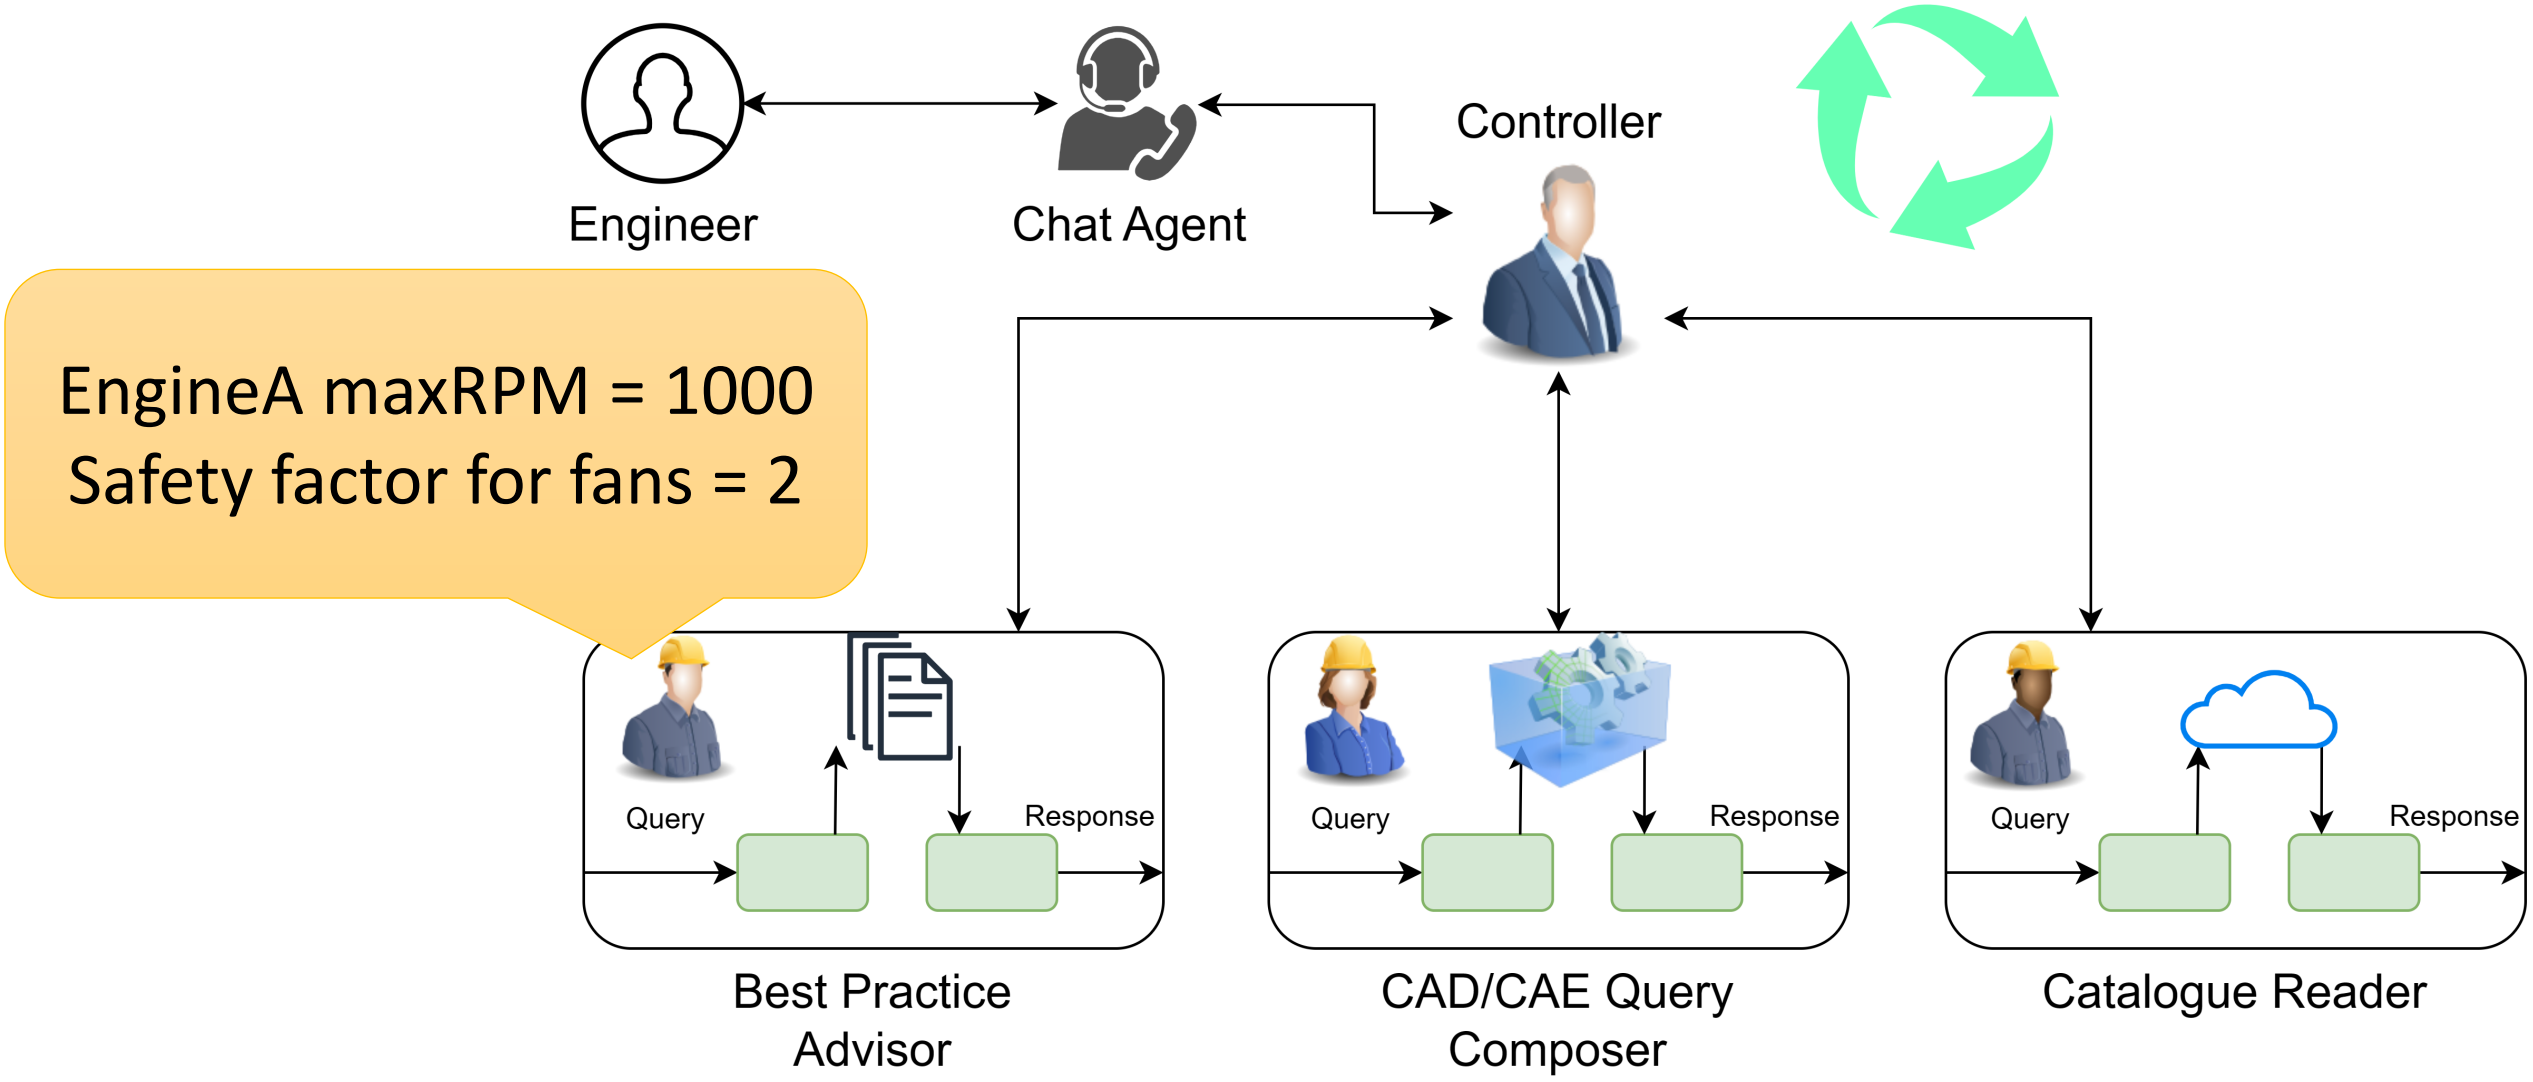

Example: Design Assistant

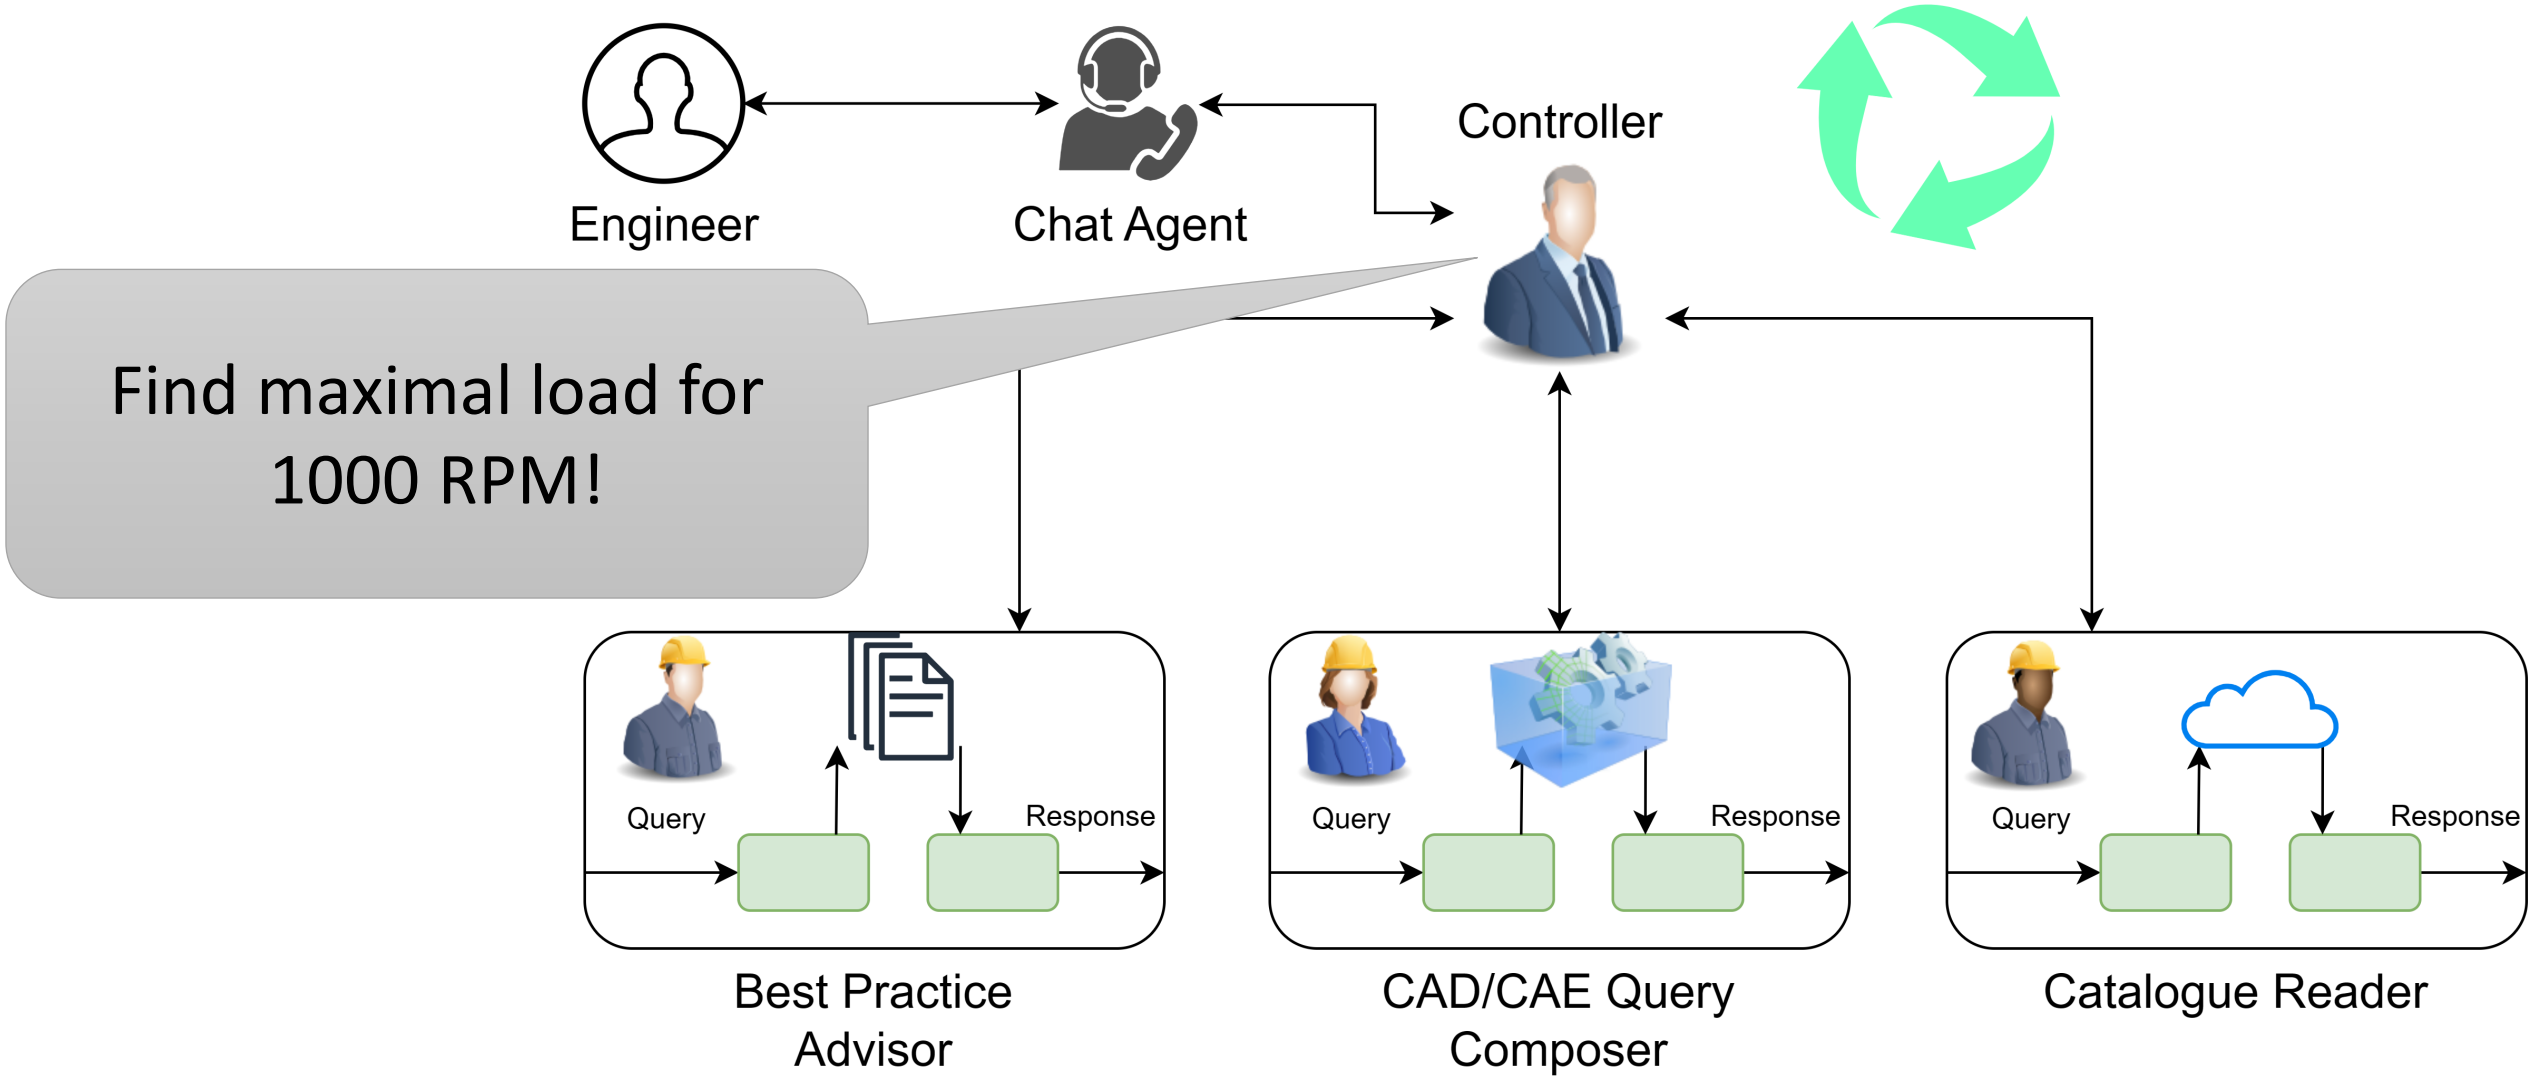

Example: Design Assistant

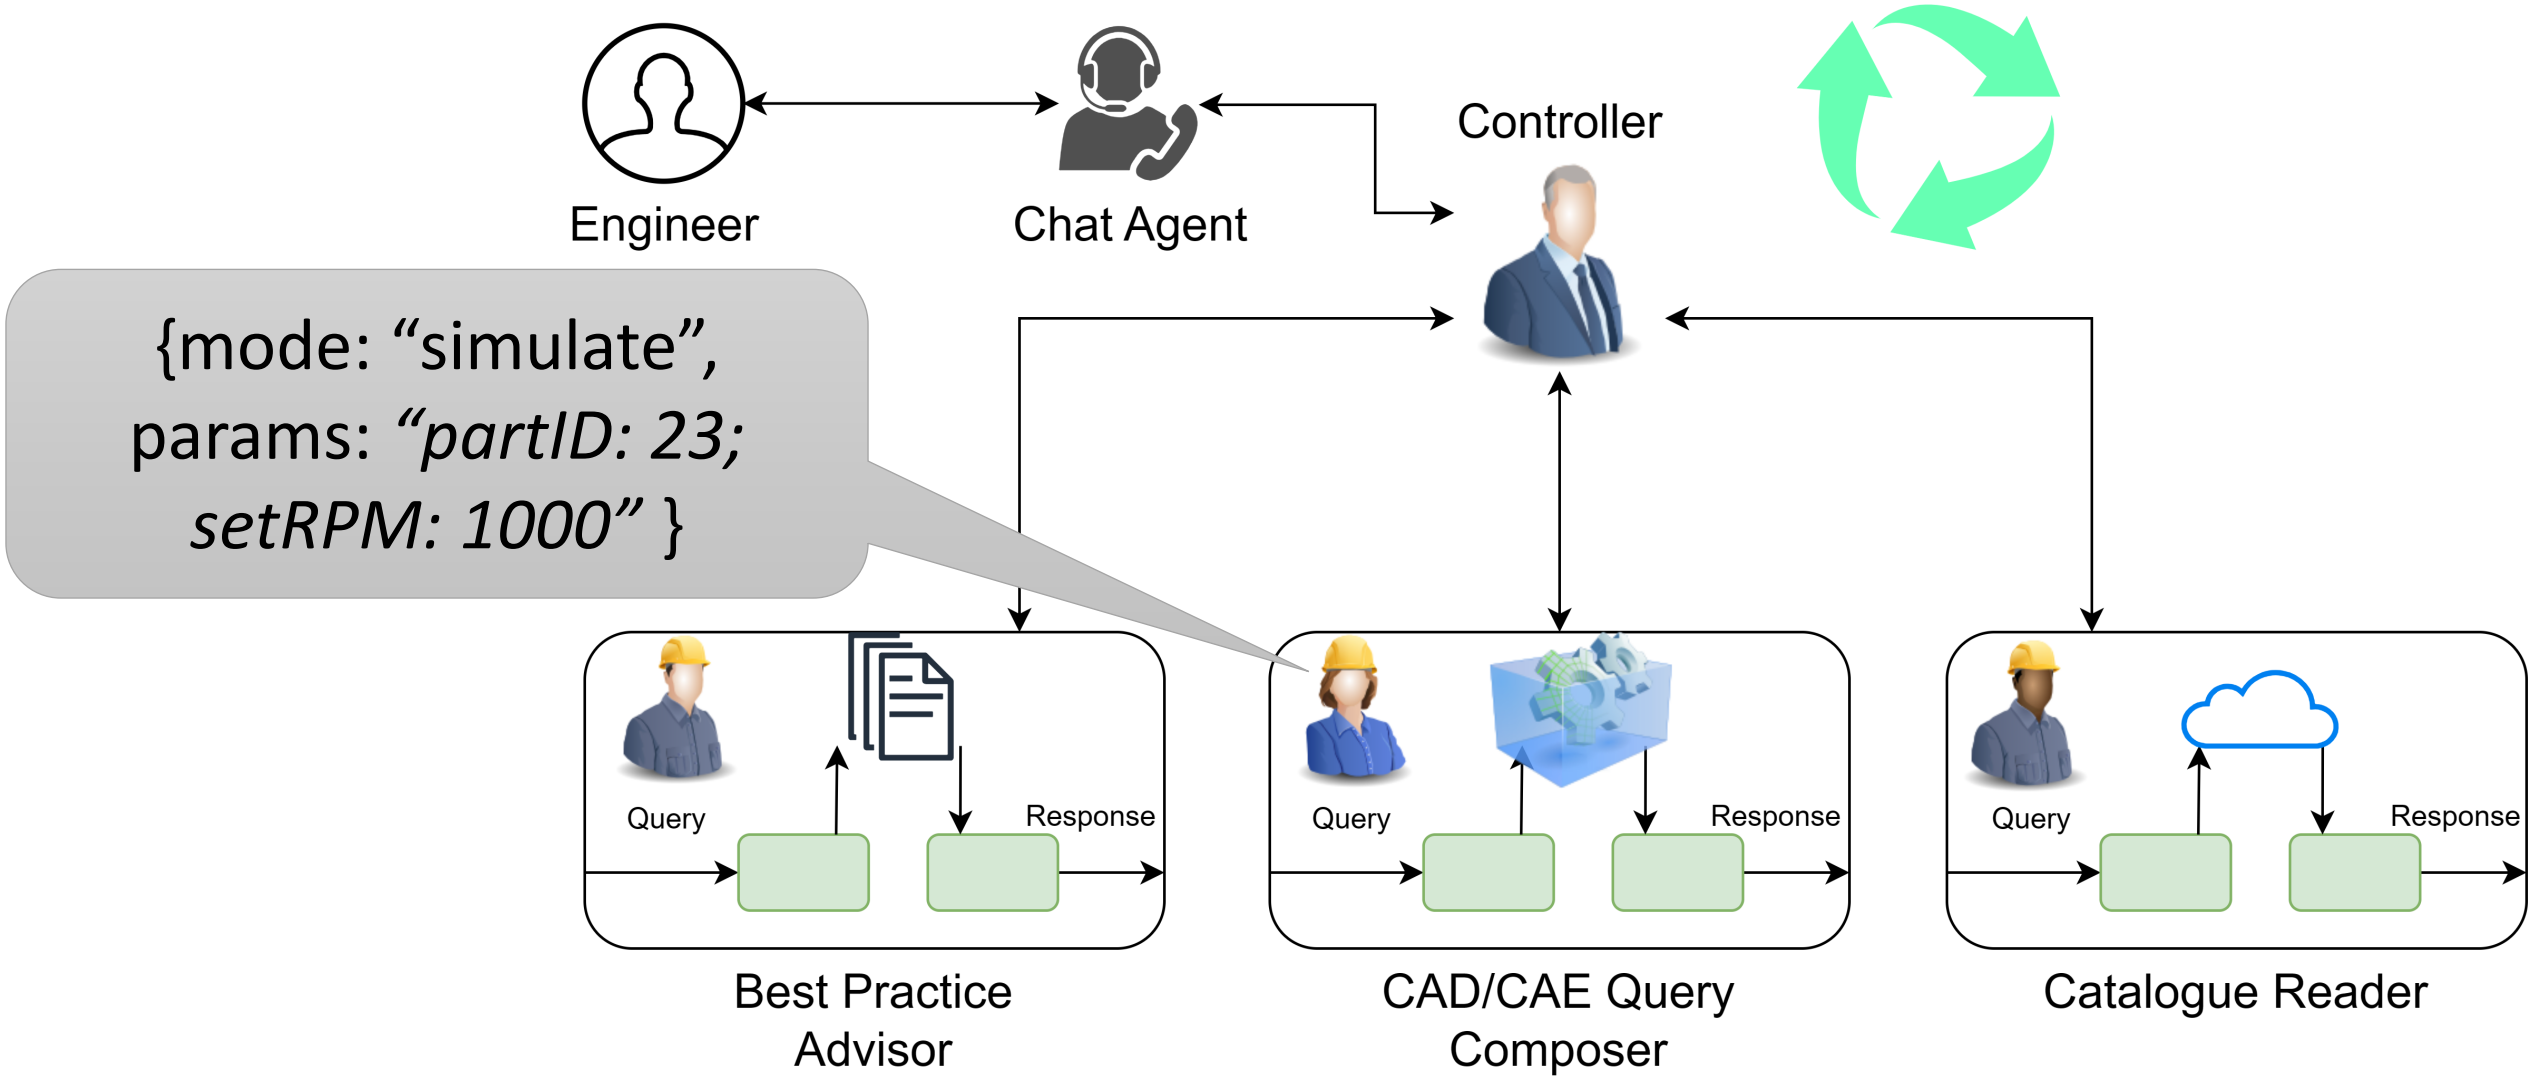

Example: Design Assistant

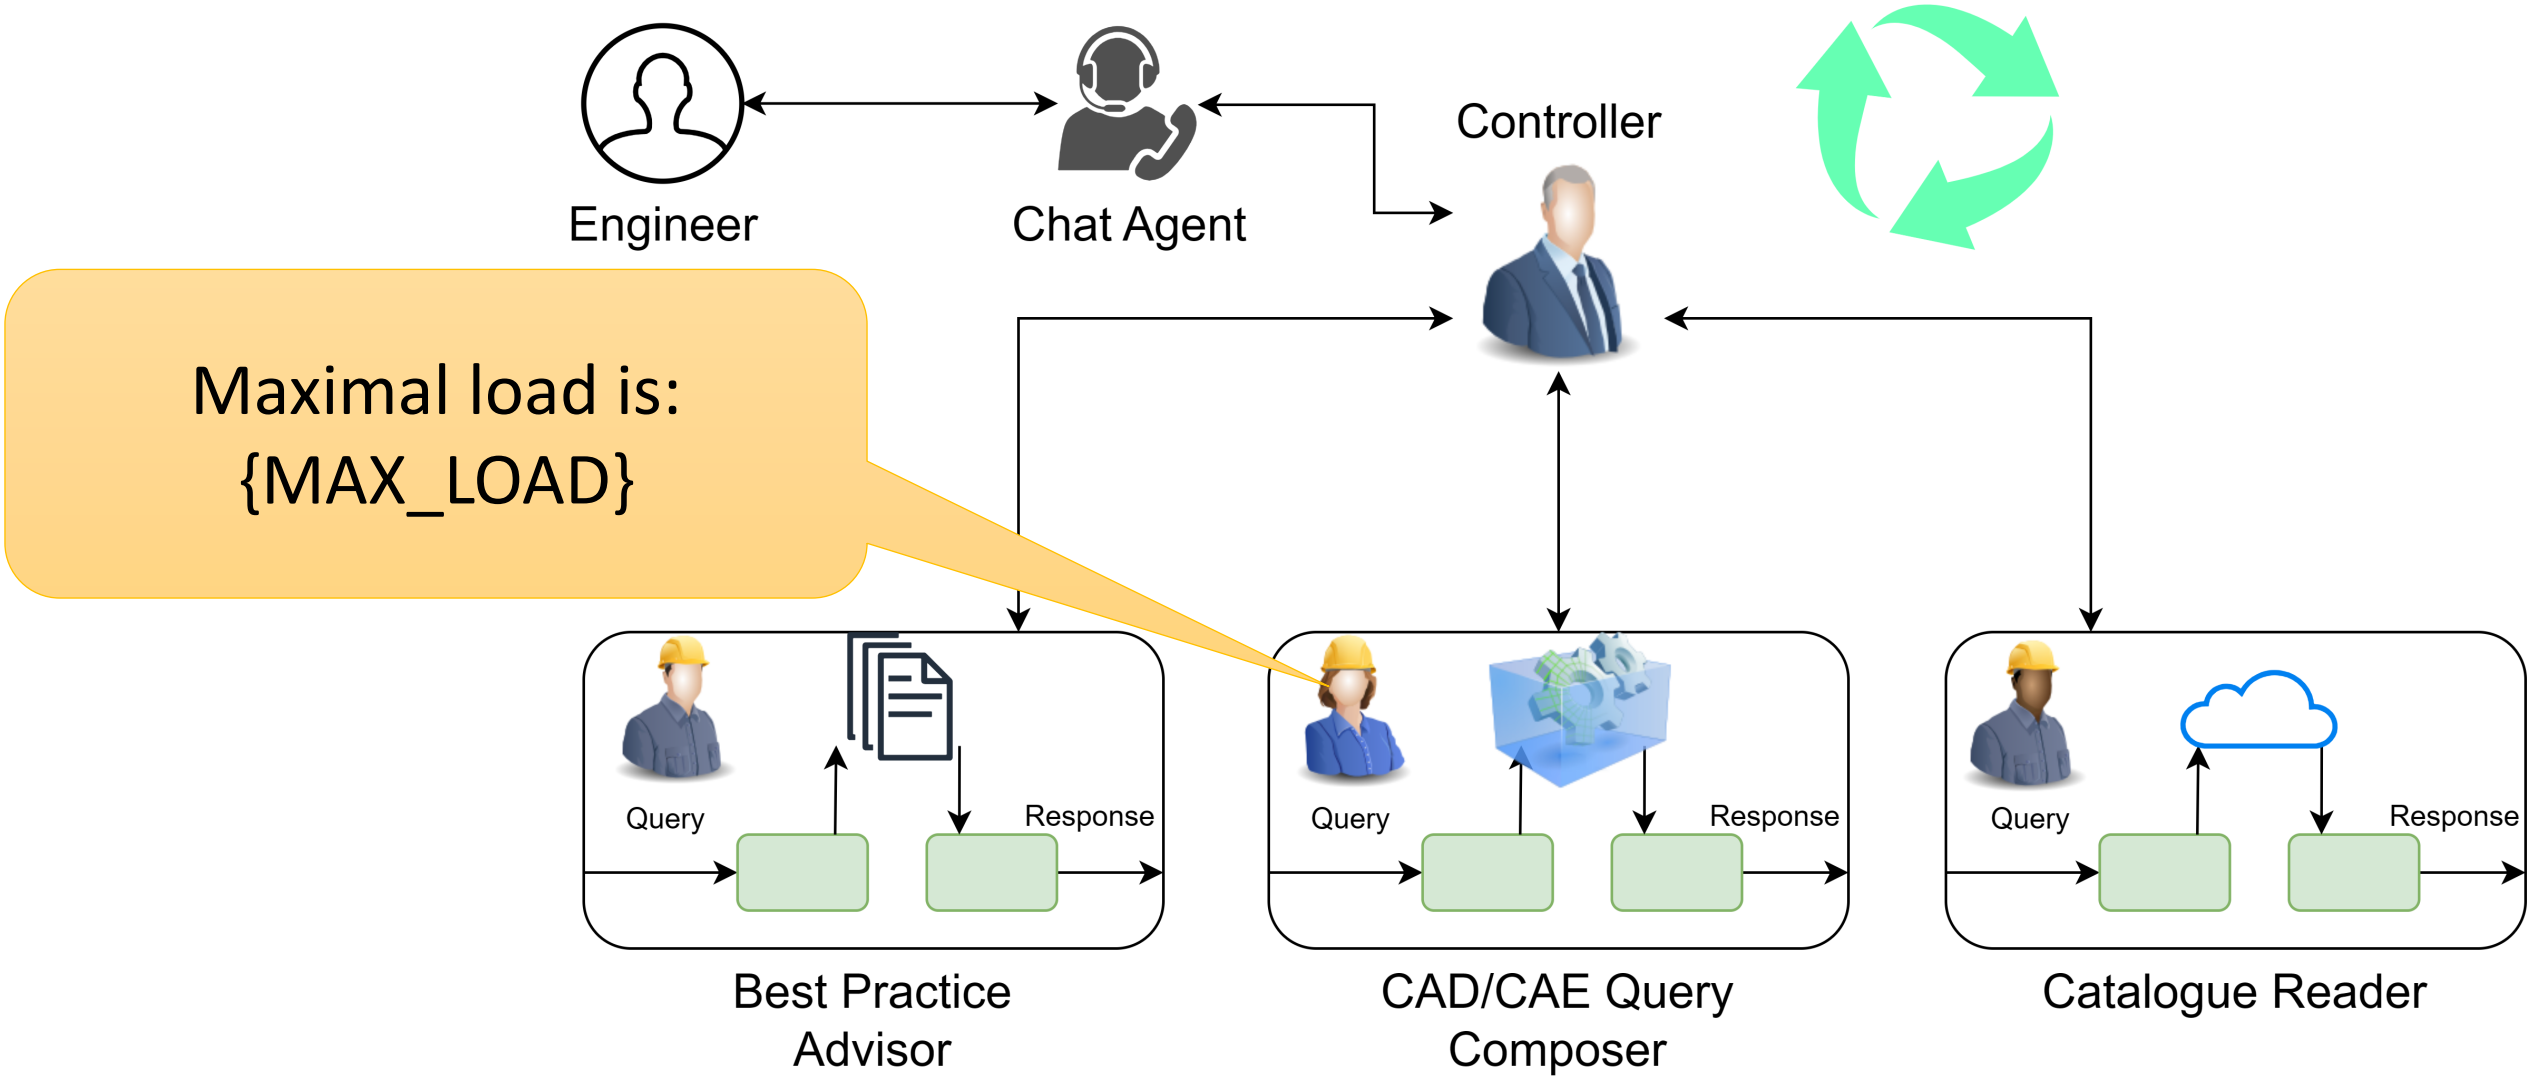

Example: Design Assistant

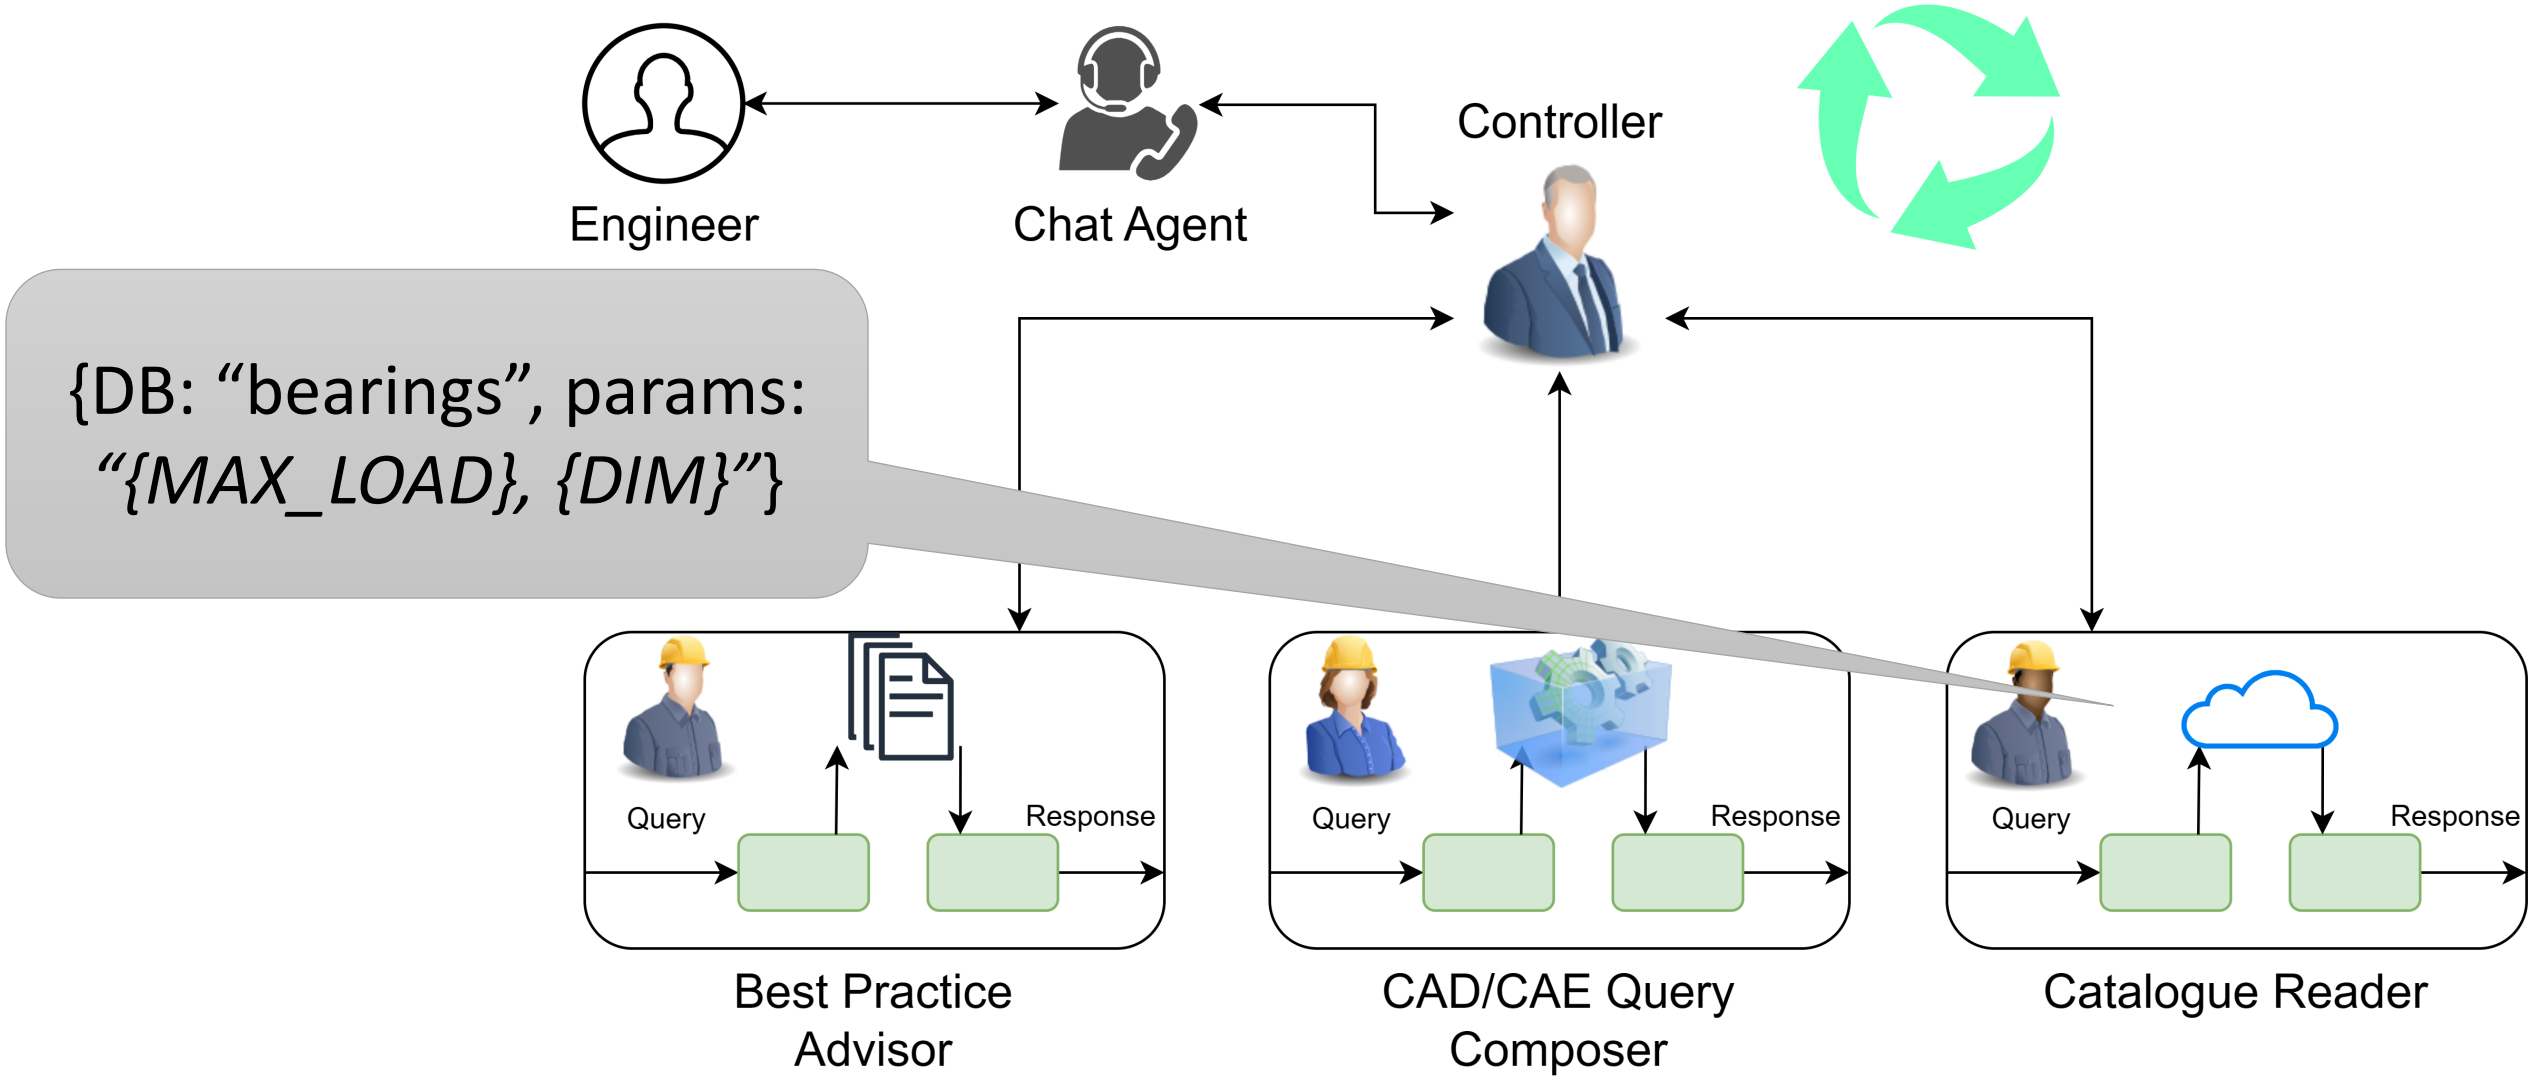

Example: Design Assistant

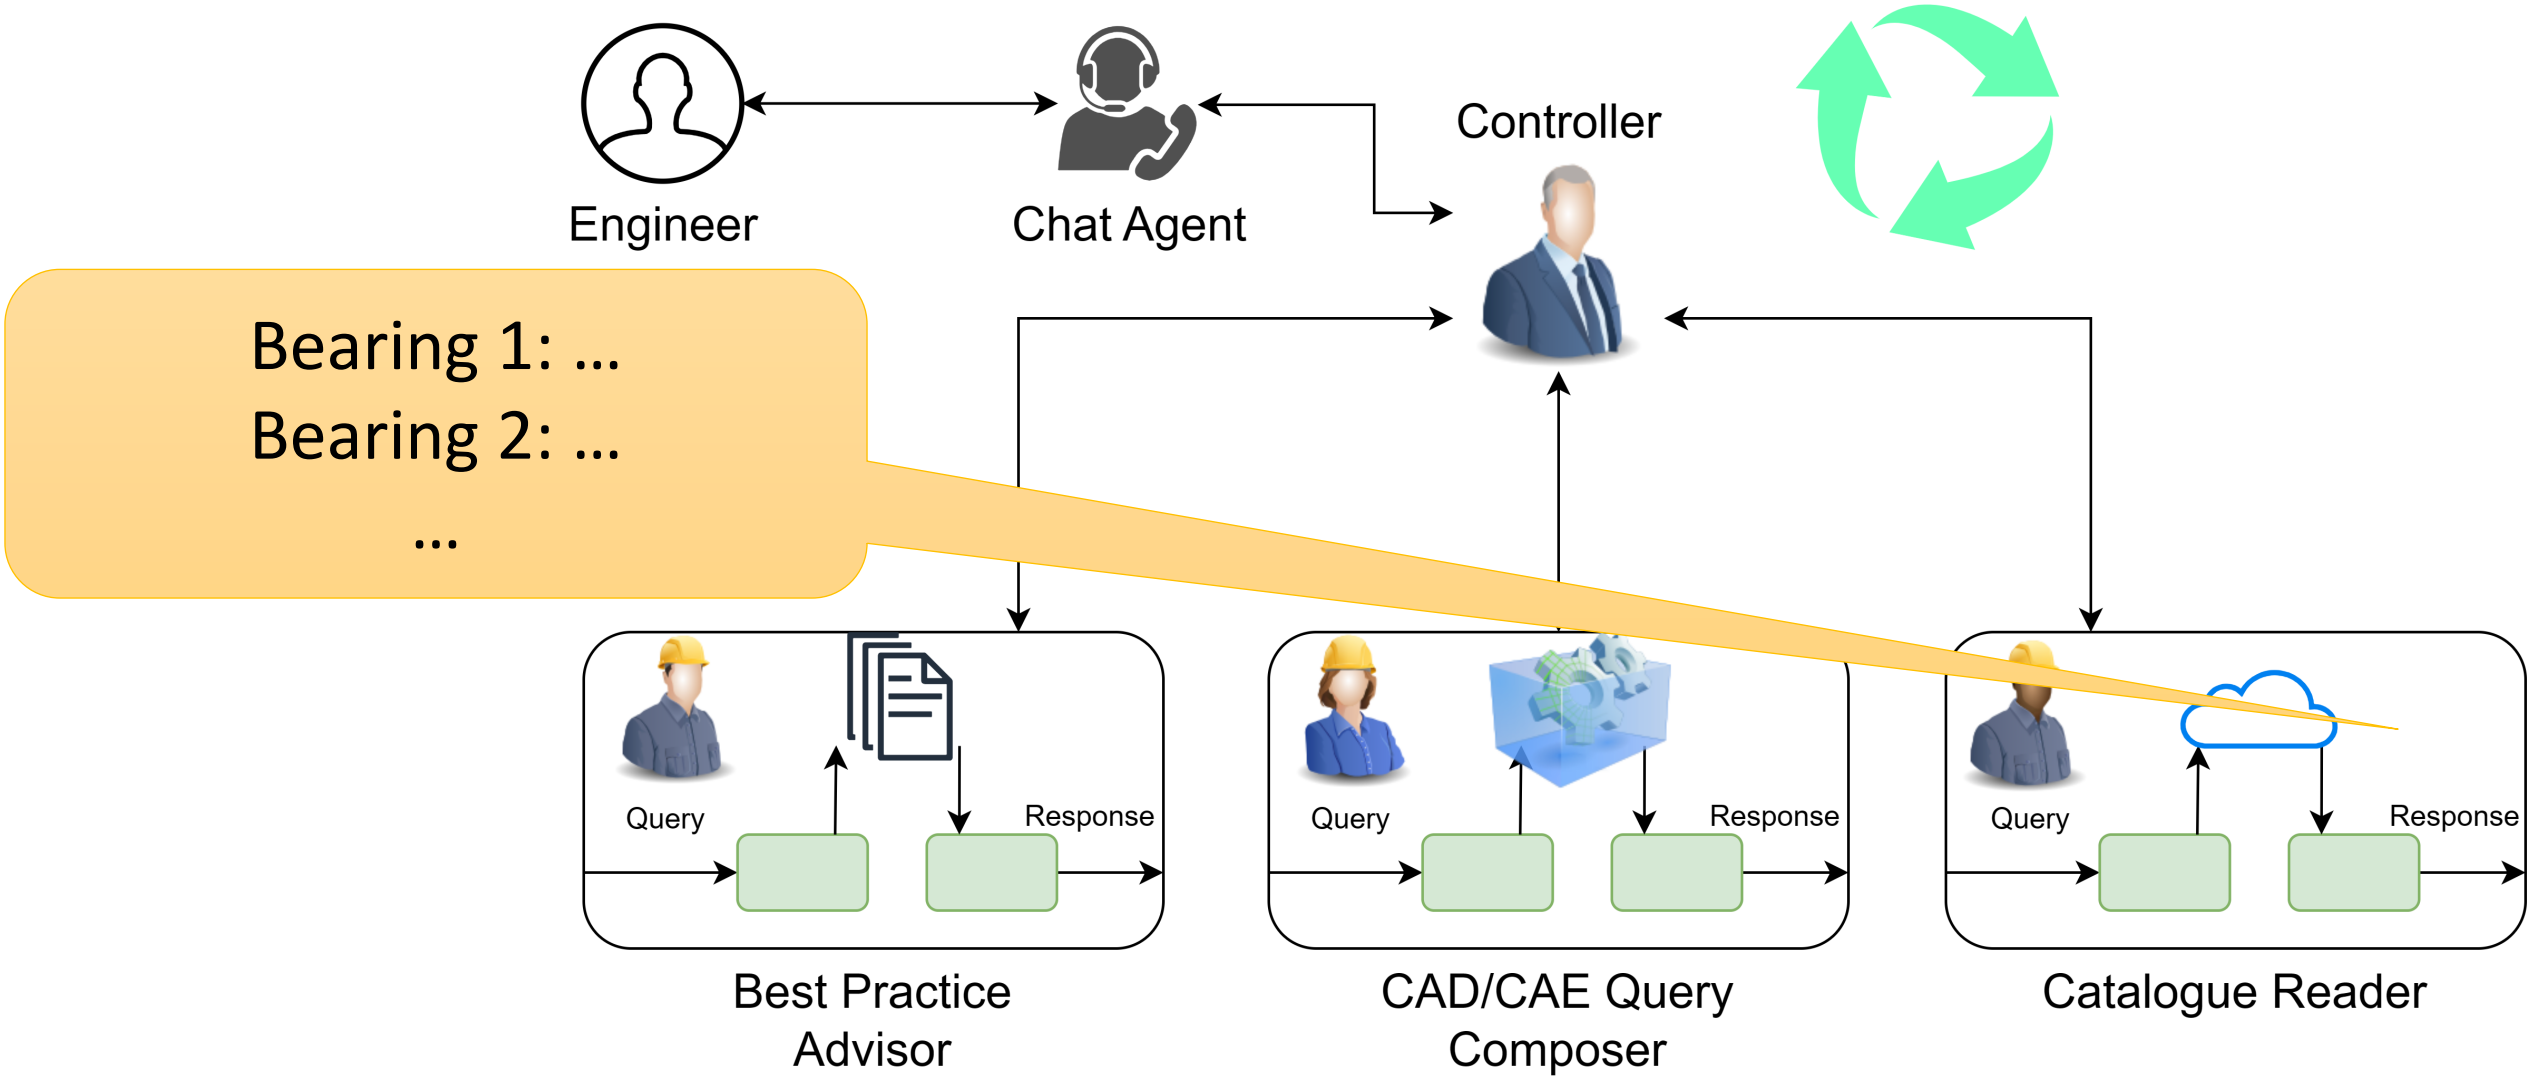

Example: Design Assistant

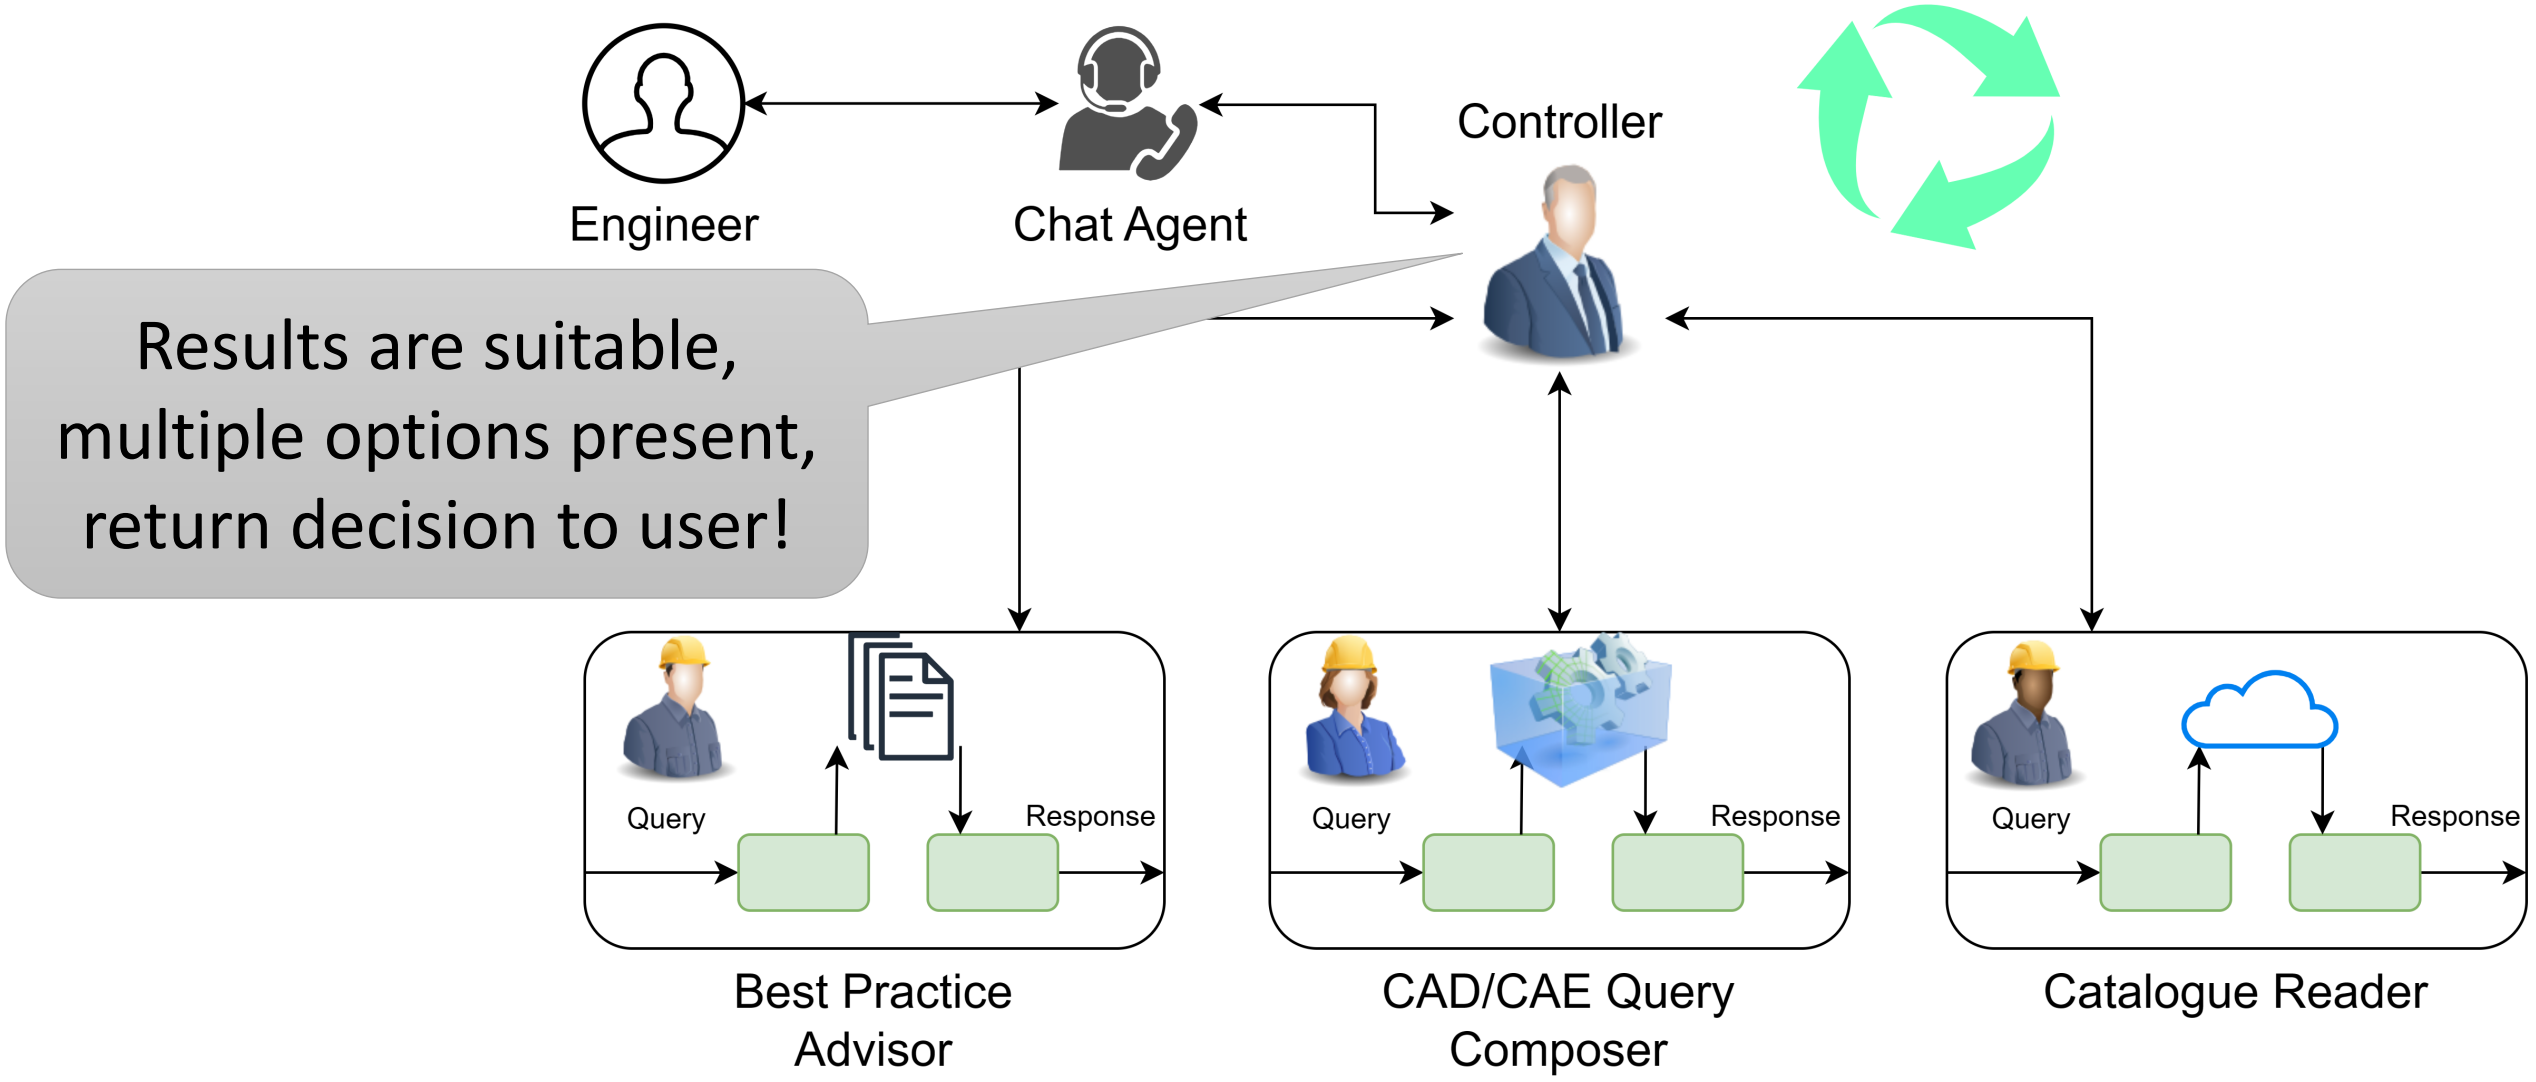

Example: Design Assistant

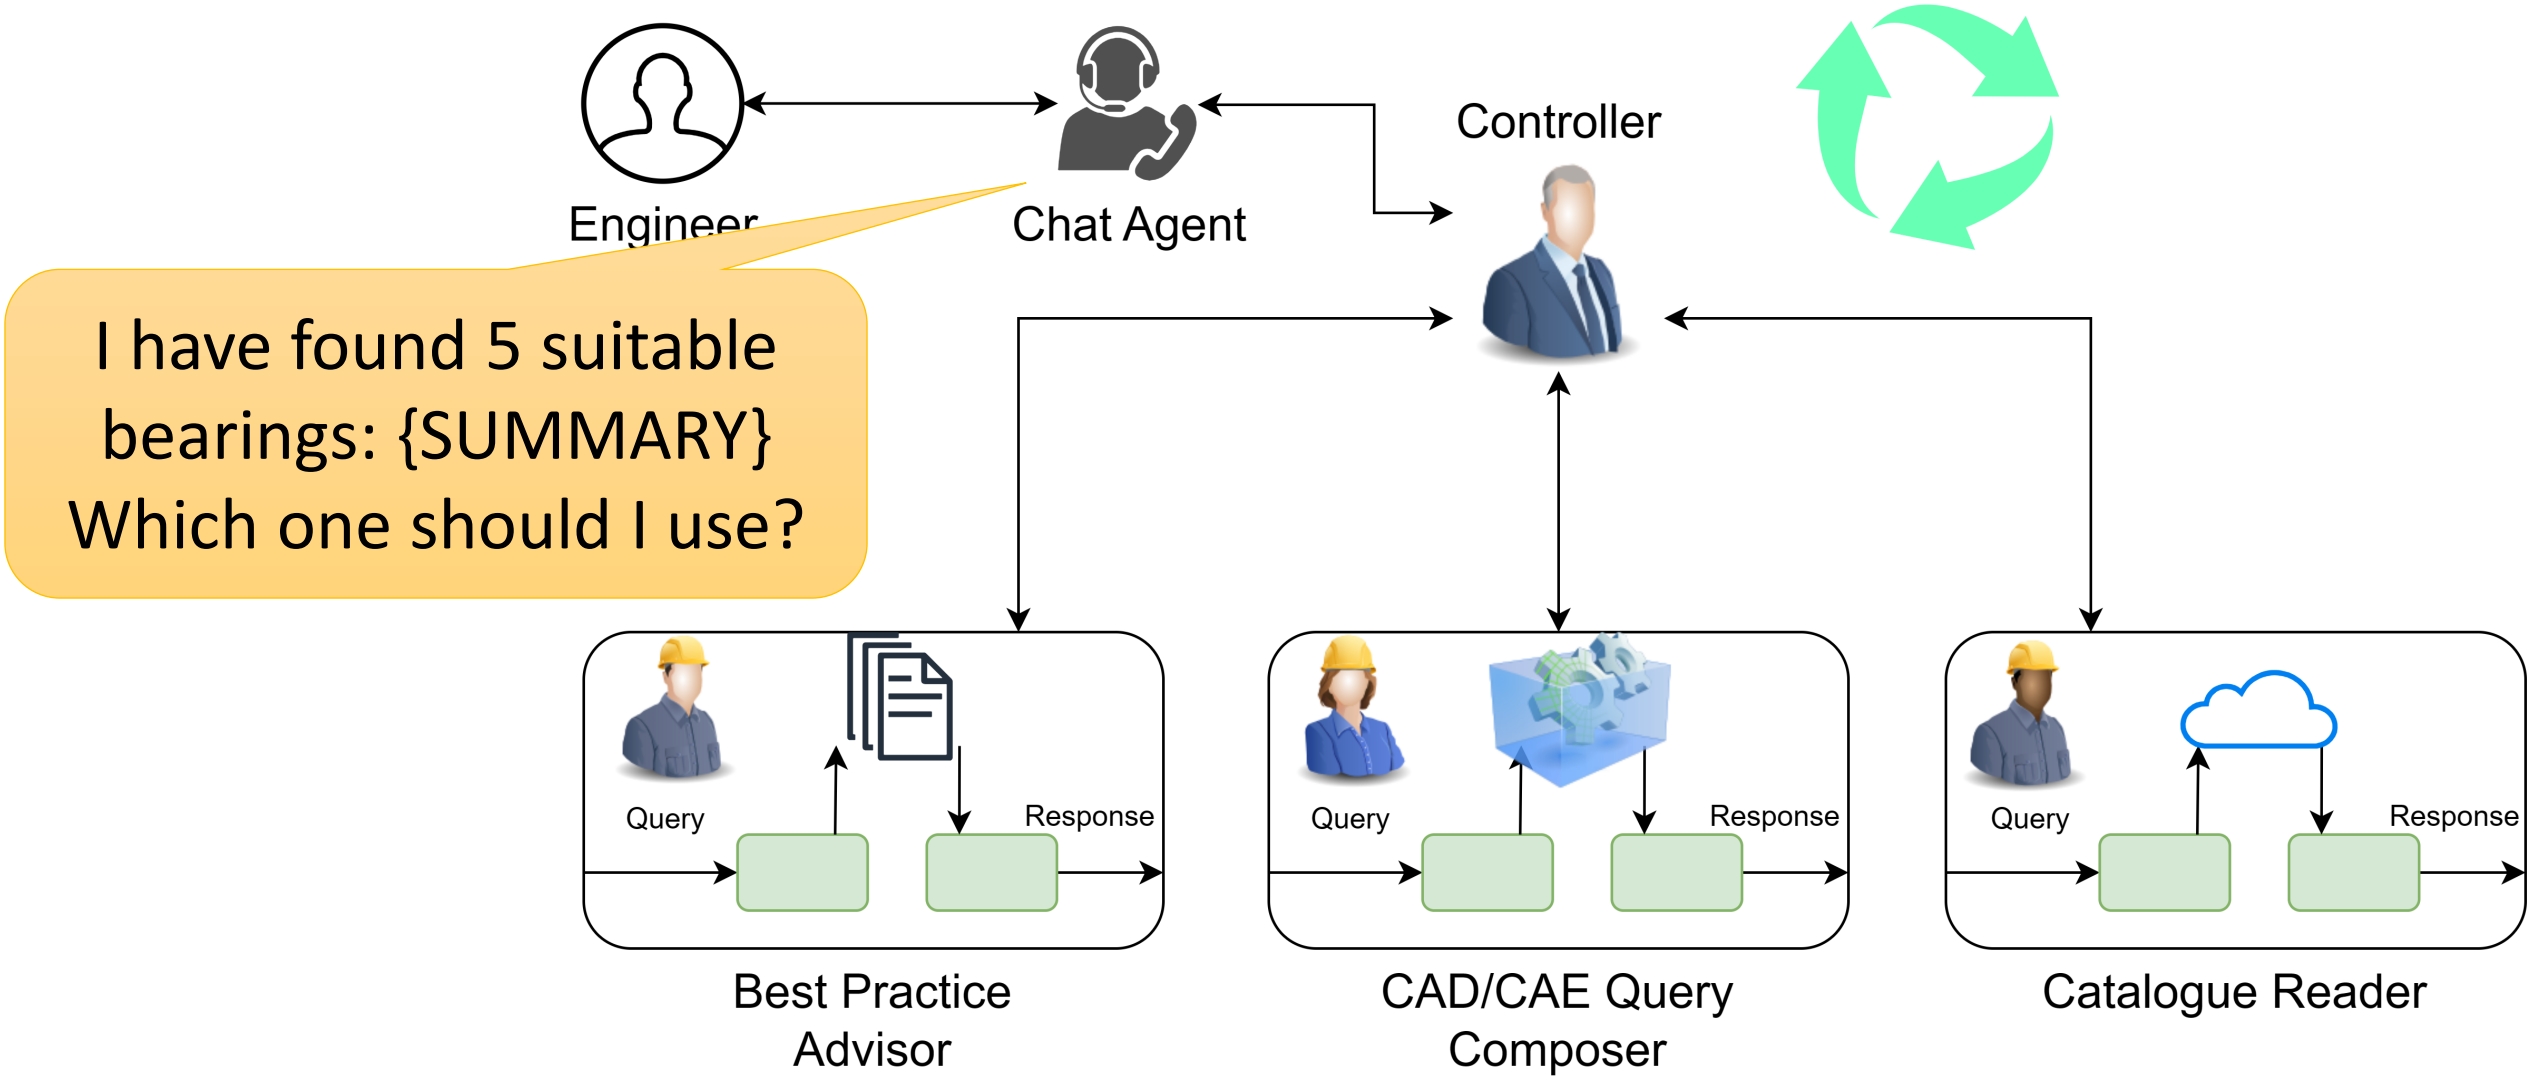

Coding Agents

- Planning included
- Tools are mostly code-related
- Might perform RAG on the repository

- Example:
- Pythagora.io

Get an OpenAI API key

- Platform.openai.com
- Dashboard
- API keys
- Create new
- Save it for later

Get the Pythagora Extension

- Install from VSCode Extensions

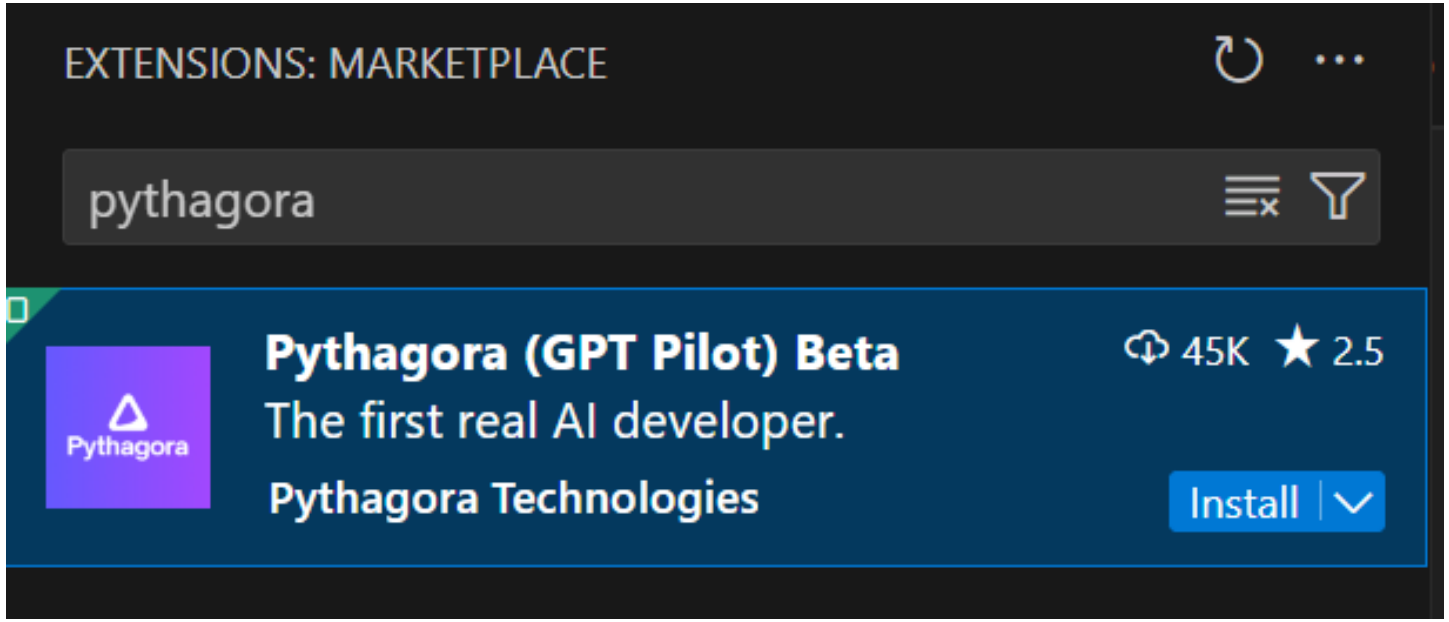

Get the Pythagora Extension

- Open the Pythagora tab, wait for it to setup
- Click on the Settings cogwheel

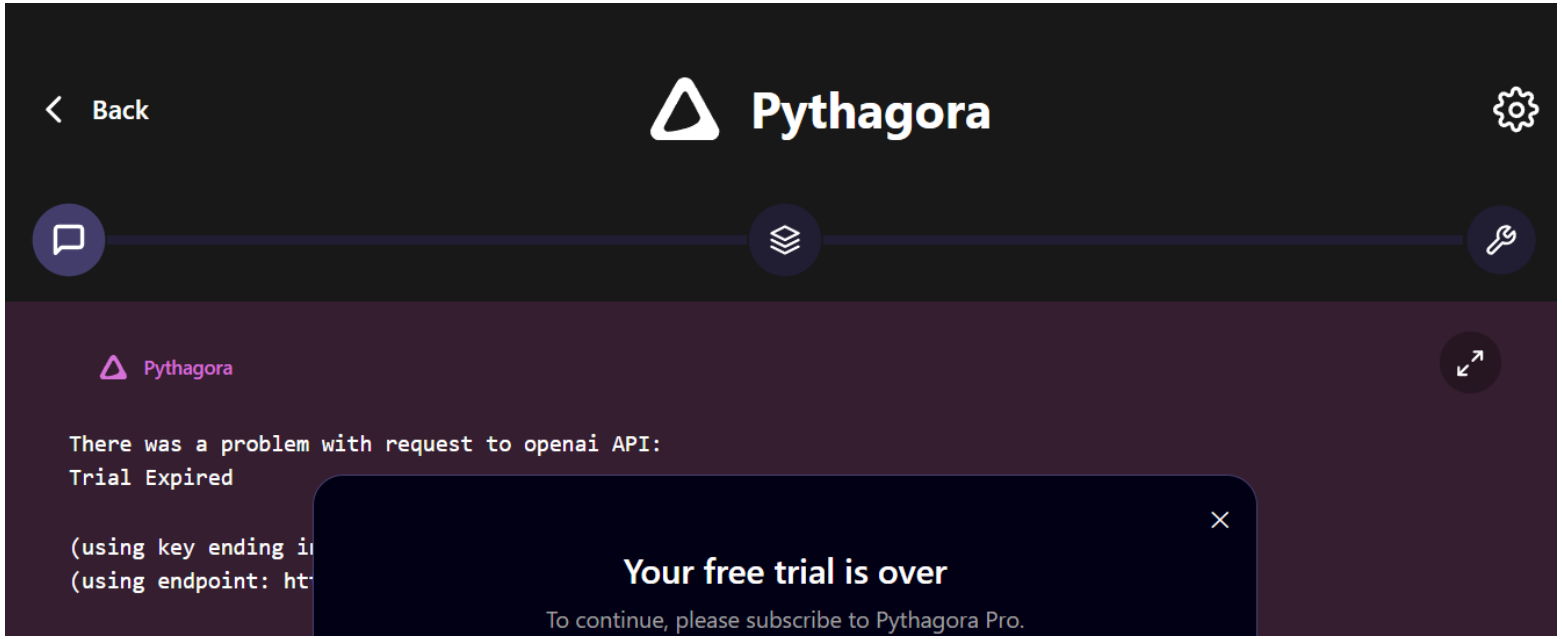

Get the Pythagora Extension

- Set and copy your GPT Pilot path to **any value**
- Navigate to that folder

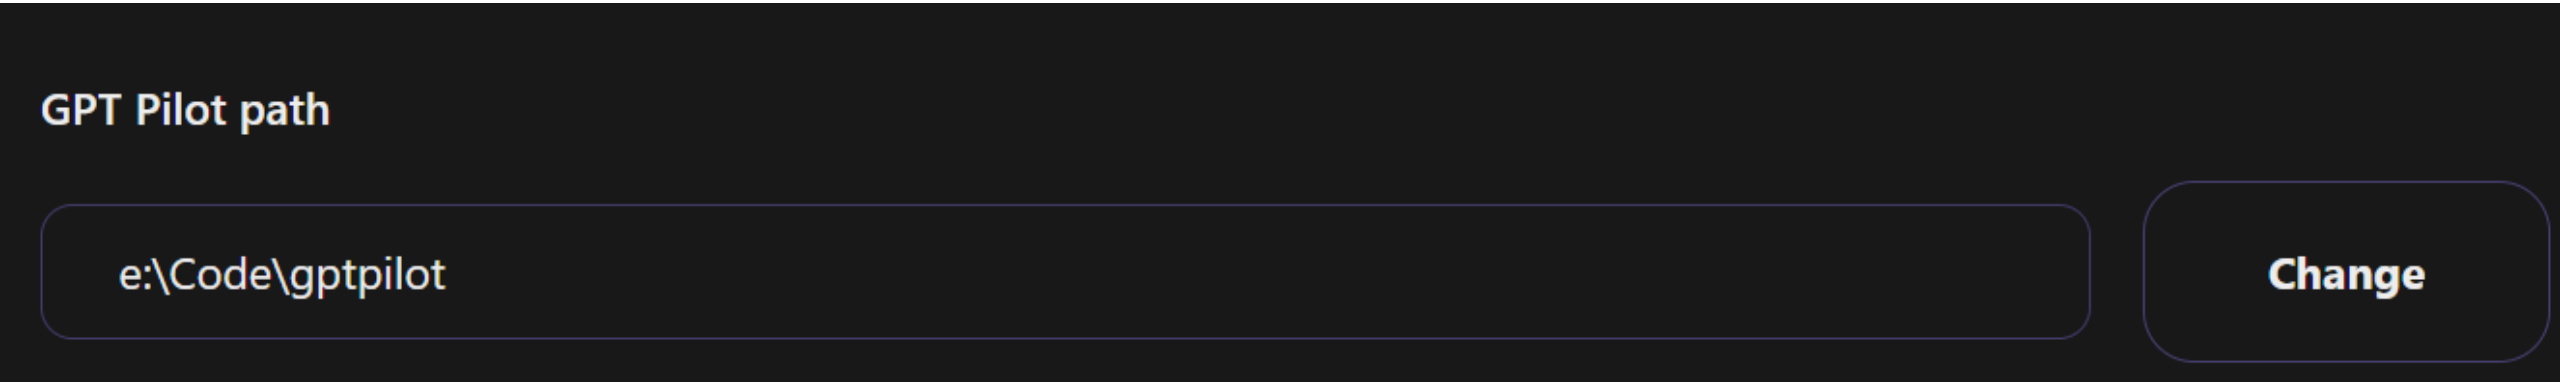

Get the Pythagora Extension

- Edit ./pilot/.env
- **OPENAI_ENDPOINT=https://api.openai.com/v1/chat/completions**
- **OPENAI_API_KEY=<YOUR KEY>**
- **MODEL_NAME=gpt-4o**

- Restart VSCode

Enjoy the interview

- Create an app answering the questions by Pythagora
- Watch it plan, act, revise!

UI Dev Tool

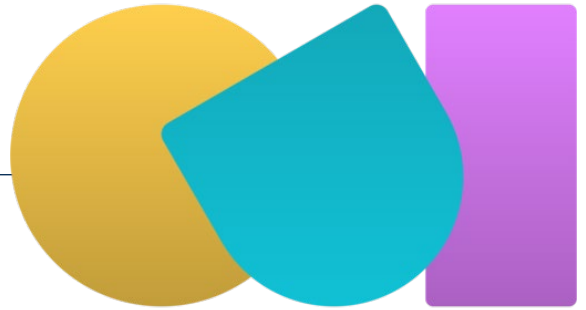

- OpenUI
- Clone from GitHub
  - git clone https://github.com/wandb/openui
  - $env:OPENAI_API_KEY = <your key>
  - cd openui/backend
  - pip install .
  - python -m openui

UI Dev Tool

- Set the model to gpt-4o in the settings
- Draw or screenshot a UI!
- Define local modifications
- Dynamic connections
- Basic animations
- etc